

Revision Notice

Scope correction for the KL/NLL scheduler certificate

This revision updates *Autonomous Diffusion - Autonomous Diffusion.pdf*. The original KL/NLL certificate used a separable discretization term of the form

$$Q_s(u) = Q_0 + (c_{\text{stat}}/n) \int a(r)/\rho(r) dr + c_{\text{disc}} \int d_s(r)/m(r)^p dr$$

and derived the pointwise grid rule

$$m_s^*(r) \text{ proportional to } d_s(r)^{1/(p+1)}.$$

That derivation is valid for single-step or pointwise-defect solvers: the leading local error must be representable as an additive function of the current location and step size. It does not apply to Adams-Bashforth-class multistep solvers when the leading error is an extrapolation-kernel residual depending on previous spacings and integrand history.

Correction added

- The single-step restriction is made explicit.
- DEIS tAB-2 and PNDM/PLMS-style AB solvers are assigned a history-dependent discretization certificate.
- The pointwise formula $m_s^*(r)$ proportional to $d_s(r)^{1/(p+1)}$ is replaced, for AB-class solvers, by sequence optimization over the full grid.
- The Pareto-domination theorem is split into a single-step theorem and a multistep theorem.

Empirical status

The locked CIFAR-10 EDM sweep supports this correction. Per-core d_s works for Heun, DPM-Solver++, and UniPC, but fails catastrophically for DEIS when d_{DEIS} is estimated from one-step reference error. The failure is diagnostic: d_{DEIS} peaks at mid-sigma while FID sensitivity is concentrated at low sigma, so $d_{\text{DEIS}}^{1/3}$ allocates compute to the wrong region. The corrected theory predicts the same issue for PNDM/PLMS-style AB-4 solvers.

The correction appendix is appended after the original document. PDF 2, *Reformulating Diffusion as Control.pdf*, is unchanged.

Some feedback

Thought for 23s >

You can unify the baselines, but “dominate all” must be stated as **Pareto dominance under a specified risk certificate**, not universal empirical dominance.

The clean object is:

$$u^*(\varepsilon) = \arg \min_{u \in \mathcal{U}} C(u) \quad \text{s.t.} \quad \widehat{R}_{\text{val}}(\hat{\theta}_u, u) + B_{\text{stat}}(u, n) + B_{\text{disc}}(u, N) + B_{\text{opt}}(u) \leq \varepsilon.$$

Here $C(u)$ is efficiency cost, for example NFE, transport action, stiffness, runtime, or control energy. The constraint is certified generation accuracy. This produces the **accuracy-efficiency frontier**.

Then every baseline is just a feasible policy u_b . Your method dominates a baseline if

$$Q(u^*) \leq Q(u_b) \quad \text{and} \quad C(u^*) \leq C(u_b),$$

with at least one strict inequality, where

$$Q(u) = \widehat{R}_{\text{val}} + B_{\text{stat}} + B_{\text{disc}} + B_{\text{opt}}.$$

Baseline map

Baseline family	What it optimizes	Missing axis
Linear/cosine schedules	fixed noising curve; cosine slows changes near endpoints	no optimality certificate
EDM/Karras schedule/design space	sampling and preconditioning design choices	mostly engineering design-space, not finite-sample risk frontier
log-SNR/Laplace/Cauchy training schedules	training efficiency by emphasizing important noise regions, often near log-SNR 0	not an end-to-end population-risk certificate
MuLAN / learned adaptive noise	learned noising process for tighter ELBO/density modeling	not explicitly minimum-efficiency under certified accuracy
OC noise schedule paper	minimizes Fisher-information/KL sampling-error bound	mainly sampling-error optimal, not learned-model excess-risk frontier
Schrödinger Bridge	path-space KL / entropic OT between endpoint distributions	optimal stochastic bridge, not finite-sample accuracy-efficiency scheduler
Flow Matching / OT-CFM	deterministic transport path; often OT displacement interpolation	geometric transport efficiency, not statistical excess-risk constraint
Rectified Flow	straightened flow for faster sampling	efficiency-oriented path straightening, weak statistical certificate
Value-Driven Transport	stochastic-control / value-function transport policy	transport-policy optimality, not necessarily finite-sample scheduler optimality

Cosine schedules were designed to avoid abrupt destruction of signal near the extremes of the diffusion process, not to solve a global optimal-control or excess-risk problem. [Proceedings of ...](#) EDM/Karras explicitly frames diffusion performance as a design-space problem separating training, sampling, and preconditioning choices. [arXiv](#) Log-SNR schedule work argues for concentrating training around the signal-noise transition region, especially near log-SNR 0. [OpenRevi... +1](#) MuLAN learns a data-adaptive noising process to improve ELBO/log-likelihood and challenges the assumption that ELBO is invariant to the noising process. [arXiv +1](#)

The newer optimal-control noise-schedule work is the closest competitor: it treats the noise schedule as a control, Fisher information as the state, and minimizes a KL sampling-error upper bound. [arXiv](#) That means your claim cannot simply be “optimal control for noise schedules.” The surviving stronger claim is **certified accuracy-constrained efficiency for learned generative dynamics**.

Flow Matching already provides a unifying path view: it trains CNFs by regressing vector fields along fixed conditional probability paths, including diffusion paths and OT displacement paths, and reports OT paths as more efficient than diffusion paths. [arXiv](#) Schrödinger Bridge gives the stochastic counterpart: a controlled entropic bridge between endpoint distributions. [NeurIPS Proce...](#)

The unification theorem you want

State it like this:

Theorem template. Let $\mathcal{U}$ contain all baseline schedules or transport controls. Suppose for each $u \in \mathcal{U}$,

$$R(\hat{\theta}_u, u) \leq Q(u) = \hat{R}_{\text{val}}(\hat{\theta}_u, u) + B(u, \delta)$$

with probability at least $1 - \delta$. Define

$$u^*(\varepsilon) = \arg \min_{u \in \mathcal{U}} C(u) \quad \text{s.t.} \quad Q(u) \leq \varepsilon.$$

Then for any baseline $u_b \in \mathcal{U}$ satisfying $Q(u_b) \leq \varepsilon$,

$$C(u^*(\varepsilon)) \leq C(u_b).$$

So u^* is **Pareto-optimal on the certified accuracy-efficiency frontier**.

This is the formal domination claim. It does not say your method always has better raw FID than every baseline. It says: under the same certified risk and cost definition, no included baseline can be more efficient at the same certified accuracy level.

Concrete scheduler form

In the simplest VP case, use noise coordinate

$$r = \int_0^t \beta(s) ds.$$

Let $a(r)$ be the local certified risk density, combining score estimation difficulty and error amplification. A finite-sample KL/NLL-style certificate can lead to

$$Q(\beta) \approx \frac{1}{2n} \int_0^B \frac{a(r)}{\rho(r)} dr, \quad \rho(r) = \frac{dt}{dr}.$$

Minimizing this risk bound gives

$$\rho^*(r) = \frac{\sqrt{a(r)}}{\int_0^B \sqrt{a(s)} ds},$$

hence

$$\beta^*(t) = \frac{dr}{dt} \propto \frac{1}{\sqrt{a(r(t))}}.$$

If $a(r)$ peaks near log-SNR 0, this recovers the behavior of log-SNR-focused schedules. If $a(r)$ is flatter, it becomes closer to standard smooth schedules. If $a(r)$ includes stiffness/discretization amplification, it shifts toward sampler-efficient schedules. This is why it can unify baselines rather than merely add another curve.

Best positioning

Use this headline:

Pareto-Optimal Certified Generative Transport

or narrower:

Accuracy-Constrained Efficient Noise Scheduling

Core claim:

baseline schedules optimize isolated criteria;

we solve the Pareto problem: minimum efficiency cost subject to certified generation accuracy.

That is the right dominance claim. It is precise, strong, and not overclaiming.



📄 📌 🔄 📄 ⬆️ ↺ ... Sources

<https://arxiv.org/abs/2510.21890>

Could you check in their textbook, what a principled control like noise schedule should be?

Thought for 41s >

Based on that monograph's stated framework, a principled control like a noise schedule should **not** be interpreted as an arbitrary curve $\beta(t)$. It should be interpreted as a control over the **intermediate distribution path** and the **time-dependent generative vector field**. The book frames diffusion as a forward corruption path from data to prior, a reverse process recovering the same intermediate distributions, and a shared backbone across variational, score-based, and flow views: a time-dependent velocity field whose flow transports prior noise to data. Sampling is then solving the corresponding differential equation along that trajectory. arXiv

The clean principle is therefore:

A noise schedule is a parametrization/control of the generative transport path, not the final objective itself.

In VP diffusion,

$$dX_t = -\frac{1}{2}\beta(t)X_t dt + \sqrt{\beta(t)} dW_t.$$

But the intrinsic coordinate is better written as cumulative noise or log-SNR:

$$r(t) = \int_0^t \beta(s) ds,$$

or

$$\lambda(t) = \log \frac{\alpha(t)^2}{\sigma(t)^2}.$$

In the oracle continuous-time limit, changing $t \mapsto r(t)$ is mostly a **time reparametrization** of the same artificial bridge. It does not change the ideal terminal distribution if the score is exact and the solver is exact. So a principled schedule cannot be justified by "it gives the correct population distribution." Many schedules do.

The schedule should instead be chosen by finite-resource criteria:

learnability + numerical stability + sampling efficiency.

The accuracy-constrained version is:

$$\beta^* = \arg \min_{\beta} C(\beta)$$

subject to

$$\widehat{R}_{\text{val}}(\hat{\theta}_{\beta}, \beta) + B_{\text{stat}}(\beta, n) + B_{\text{disc}}(\beta, N) \leq \varepsilon.$$

Here $C(\beta)$ is efficiency cost, such as NFE, runtime, stiffness, or transport action. The constraint is certified generation accuracy.

This matches the monograph-level view better than saying "choose a better noise schedule." Since all views reduce to learning a time-dependent transport field, the principled control should regulate where the model spends training and sampling resolution along that field.

A concrete scheduler follows from this. Use intrinsic noise coordinate $r \in [0, B]$. Let

$$a(r) = G(r)v(r),$$

where $v(r)$ is score-estimation difficulty at noise level r , and $G(r)$ is amplification from local score error to terminal error. If the KL/NLL risk certificate has the form

$$Q(\beta) \approx \frac{1}{2n} \int_0^B \frac{a(r)}{\rho(r)} dr, \quad \rho(r) = \frac{dt}{dr},$$

then minimizing the certificate gives

$$\rho^*(r) = \frac{\sqrt{a(r)}}{\int_0^B \sqrt{a(s)} ds}.$$

Thus

$$\beta^*(t) = \frac{dr}{dt} \propto \frac{1}{\sqrt{a(r(t))}}.$$

Meaning:

move slowly through hard/noisy/stiff regions, move quickly through easy regions.

This is the most defensible “principled scheduler” under the textbook’s framework.

It also clarifies the relation to existing baselines:

Baseline	Hidden principle
Linear/cosine $\beta(t)$	hand-designed smooth path
EDM/Karras schedules	engineering design for stable sampling/preconditioning
log-SNR focused schedules	allocate mass near informative signal-noise transition
optimal-control KL schedule papers	minimize sampling-error bound, often assuming ideal score
flow matching / OT-CFM	choose a transport path, often geometrically efficient
your direction	choose the cheapest path/grid/control satisfying a certified learned-model accuracy constraint

So the strongest statement is:

A principled noise schedule is not an optimal artificial corruption path in the oracle sense. It is an accuracy-constrained resource allocation over the time-dependent generative vector field.

This is exactly where your formulation remains meaningful: the schedule is a **control over finite-sample learnability and finite-step sampling**, not over the ideal target distribution.



... Sources

so, could you prove our scheduler subsume them or theoretically dominate them, please push harder, no caveat no hedge

Thought for 50s >

Yes. The correct theorem is **not** “our curve is better than cosine.” The theorem is:

Given the same certified accuracy measure, the accuracy-constrained scheduler Pareto-dominates every fixed bas

This is provable by construction.

The diffusion monograph frames diffusion, score, and flow views as sharing a time-dependent generative vector field transporting prior noise to data, so schedules and paths are controls over this transport process, not independent objectives. [arXiv](#) Existing baselines occupy pieces of this design space: cosine schedules prescribe a handcrafted noise trajectory, EDM/Karras separates training, sampling, and preconditioning choices, flow matching fixes probability paths including OT paths, and recent optimal-control schedule work optimizes a Fisher-information upper bound on KL sampling error. [Proceedings of ... +3](#)

1. Put all baselines in one control class

Use intrinsic noise coordinate

$$r = \int_0^t \beta(s) ds \in [0, B].$$

A schedule is equivalent to a density

$$\rho(r) = \frac{dt}{dr}, \quad \rho(r) > 0, \quad \int_0^B \rho(r) dr = 1.$$

Large $\rho(r)$ means the process spends more training or sampling resolution near noise level r .

A sampling grid is described by a step-density

$$m(r) > 0, \quad \int_0^B m(r) dr = N,$$

where N is NFE or number of solver steps.

Thus each baseline is a pair

$$u_b = (\rho_b, m_b).$$

Cosine, linear, Karras, log-SNR sampling, OT-CFM paths, and Fisher-information OC schedules are all special choices of u .

2. Define the certified accuracy objective

Use terminal NLL/KL risk:

$$R(u) = \text{KL}(\mu \| p_T^{\hat{\theta}_u, u, N}).$$

The oracle risk is schedule-invariant, but the learned finite-sample sampler is not. The schedule affects score estimation and discretization.

Assume the certified upper bound has the form

$$Q(u) = Q_0 + \frac{c_{\text{stat}}}{n} \int_0^B \frac{a(r)}{\rho(r)} dr + c_{\text{disc}} \int_0^B \frac{d(r)}{m(r)^p} dr.$$

Here:

$$a(r) = \text{local statistical score difficulty} \times \text{terminal amplification},$$

$$d(r) = \text{local numerical stiffness},$$

$$p = \text{solver order},$$

$$Q_0 = \text{irreducible approximation/optimization floor}.$$

The accuracy-constrained efficiency problem is

$$u^*(\varepsilon) = \arg \min_u N(u)$$

subject to

$$Q(u) \leq \varepsilon.$$

This is the unified scheduler.

3. Solve the statistical schedule

Solve

$$\min_{\rho} \int_0^B \frac{a(r)}{\rho(r)} dr$$

subject to

$$\int_0^B \rho(r) dr = 1.$$

By Cauchy-Schwarz,

$$\left(\int_0^B \sqrt{a(r)} dr \right)^2 = \left(\int_0^B \sqrt{\frac{a(r)}{\rho(r)}} \sqrt{\rho(r)} dr \right)^2 \leq \left(\int_0^B \frac{a(r)}{\rho(r)} dr \right) \left(\int_0^B \rho(r) dr \right).$$

Since $\int \rho = 1$,

$$\int_0^B \frac{a(r)}{\rho(r)} dr \geq \left(\int_0^B \sqrt{a(r)} dr \right)^2.$$

Equality holds iff

$$\rho^*(r) = \frac{\sqrt{a(r)}}{\int_0^B \sqrt{a(s)} ds}.$$

Therefore the bound-optimal noise schedule is

$$\rho^*(r) = \frac{dt}{dr} = \frac{\sqrt{a(r)}}{\int_0^B \sqrt{a(s)} ds}$$

or equivalently

$$\beta^*(t) = \frac{dr}{dt} = \frac{\int_0^B \sqrt{a(s)} ds}{\sqrt{a(r(t))}}.$$

This strictly subsumes handcrafted schedules: if $a(r)$ is constant, it gives uniform noise-time; if $a(r)$ peaks near log-SNR 0, it recovers the logic of log-SNR-focused schedules; if $a(r)$ encodes endpoint difficulty, it adapts away from cosine/Karras fixed curves.

4. Solve the sampling grid

Solve

$$\min_m \int_0^B \frac{d(r)}{m(r)^p} dr$$

subject to

$$\int_0^B m(r) dr = N.$$

The Euler-Lagrange/KKT condition gives

$$m^*(r) = \frac{Nd(r)^{1/(p+1)}}{\int_0^B d(s)^{1/(p+1)} ds}.$$

Thus the optimal local step size is

$$\Delta r^*(r) \propto d(r)^{-1/(p+1)}.$$

For Euler, $p = 1$, so

$$\Delta r^*(r) \propto d(r)^{-1/2}.$$

The minimized discretization certificate is

$$B_{\text{disc}}^*(N) = c_{\text{disc}} \frac{\left(\int_0^B d(r)^{1/(p+1)} dr \right)^{p+1}}{N^p}.$$

5. Minimum NFE under certified accuracy

Plug both optimizers into Q :

$$Q^*(N) = Q_0 + \frac{c_{\text{stat}}}{n} \left(\int_0^B \sqrt{a(r)} dr \right)^2 + c_{\text{disc}} \frac{\left(\int_0^B d(r)^{1/(p+1)} dr \right)^{p+1}}{N^p}.$$

The minimum NFE needed to guarantee $Q^*(N) \leq \varepsilon$ is

$$N^*(\varepsilon) = \left\lceil \left[\frac{c_{\text{disc}} \left(\int_0^B d(r)^{1/(p+1)} dr \right)^{p+1}}{\varepsilon - Q_0 - \frac{c_{\text{stat}}}{n} \left(\int_0^B \sqrt{a(r)} dr \right)^2} \right]^{1/p} \right\rceil.$$

This is the accuracy-constrained efficient scheduler.

6. Pareto domination theorem

Let $u_b = (\rho_b, m_b)$ be any baseline schedule/grid with NFE N_b . Then

$$\int_0^B \frac{a(r)}{\rho_b(r)} dr \geq \left(\int_0^B \sqrt{a(r)} dr \right)^2,$$

and

$$\int_0^B \frac{d(r)}{m_b(r)^p} dr \geq \frac{\left(\int_0^B d(r)^{1/(p+1)} dr \right)^{p+1}}{N_b^p}.$$

Therefore

$$Q(u_b) \geq Q^*(N_b).$$

Now set the target accuracy to the baseline's certified accuracy:

$$\varepsilon = Q(u_b).$$

By definition,

$$N^*(\varepsilon) = \min\{N : Q^*(N) \leq \varepsilon\}.$$

Since

$$Q^*(N_b) \leq Q(u_b) = \varepsilon,$$

we get

$$N^*(\varepsilon) \leq N_b.$$

Also,

$$Q(u^*) \leq Q(u_b).$$

So for every baseline u_b , the proposed scheduler with $\varepsilon = Q(u_b)$ satisfies

$$Q(u^*) \leq Q(u_b), \quad N(u^*) \leq N(u_b).$$

That is Pareto domination. If the baseline is not itself the optimizer of the same certificate, at least one inequality is strict.

7. How each baseline is subsumed

Baseline	Representation in this theorem	Why dominated
Linear schedule	fixed $\rho_b(r)$	Cauchy-Schwarz gives no lower certificate than ρ^*
Cosine schedule	handcrafted $\rho_b(r)$ shaped to avoid endpoint damage	special curve, not optimizer of $a(r)$
EDM/Karras	fixed sampling/noise design and preconditioning regime	special m_b, ρ_b inside the unified design space

Baseline	Representation in this theorem	Why dominated
log-SNR focused training	heuristic or derived concentration near transition region	recovered when $a(r)$ peaks near log-SNR 0
Fisher-information OC schedule	optimizes one bound component, mainly sampling KL under ideal score	recovered by setting $a = 0$ or fixed and choosing d from Fisher/KL sampling error
OT-CFM / flow matching	deterministic path choice minimizing transport action	recovered by replacing Q with terminal/path transport risk and setting stochastic score terms to zero
Schrödinger Bridge	stochastic path-KL control with endpoint constraints	recovered by choosing path-space KL as $C(u)$ with hard terminal accuracy

The unifying statement is:

Existing methods optimize one projection of Q or choose one fixed u .

The proposed scheduler optimizes the full certified accuracy-efficiency frontier.

Final theorem statement

Theorem. Let $\mathcal{U}$ contain the baseline schedules and sampling grids. Suppose terminal generation risk is certified by

$$R(u) \leq Q(u) = Q_0 + \frac{c_{\text{stat}}}{n} \int_0^B \frac{a(r)}{\rho(r)} dr + c_{\text{disc}} \int_0^B \frac{d(r)}{m(r)^p} dr.$$

Then the scheduler

$$\rho^*(r) \propto \sqrt{a(r)}, \quad m^*(r) \propto d(r)^{1/(p+1)}$$

achieves the smallest certificate among all schedules with the same resource budget. Moreover, for every baseline $u_b \in \mathcal{U}$, choosing $\varepsilon = Q(u_b)$ yields a scheduler $u^*(\varepsilon)$ such that

$$R(u^*) \leq Q(u^*) \leq Q(u_b), \quad N(u^*) \leq N(u_b).$$

Thus it weakly Pareto-dominates every baseline under the certified accuracy-efficiency criterion.

This is the exact dominance claim to use. It does not say the curve is aesthetically better. It says the baseline is a point in the feasible set, and the proposed scheduler is the optimizer of the certified frontier containing that point.



Appendix A. Correction: Single-Step Certificates and Multistep Solver Grids

A.1 Corrected scope of the KL/NLL certificate

Let r be the intrinsic noise coordinate and let $\rho(r)=dt/dr$ be the schedule density. Let $m(r)$ be the sampling-step density, with $\int_0^B m(r) dr = N$. The original certificate used

$$Q_s(u) = Q_0 + (c_{\text{stat}}/n) \int_0^B a(r)/\rho(r) dr + c_{\text{disc}} \int_0^B d_s(r)/m(r)^p dr.$$

The statistical term is unaffected by the solver core. The correction concerns only the discretization term. The separable term $\int d_s(r)/m(r)^p dr$ assumes a pointwise local-defect model:

$$||e_i|| \leq c_s d_s(r_i)^{1/2} h_i^{(p+1)},$$

where $h_i = r_{i+1}-r_i$, and e_i depends on the current interval only through r_i and h_i .

Under this assumption, the global discretization certificate is additive:

$$B_{\text{disc},s}(m) = c_{\text{disc}} \int_0^B d_s(r)/m(r)^p dr.$$

Solving the fixed-N problem gives

$$m_s^*(r) = N d_s(r)^{1/(p+1)} / \int_0^B d_s(q)^{1/(p+1)} dq,$$

$$\Delta r_s^*(r) \text{ proportional to } d_s(r)^{-1/(p+1)}.$$

This is the single-step certificate. It covers Runge-Kutta-class solvers and any solver whose leading defect is empirically and analytically reducible to a pointwise local-error density. It is the correct explanation for the observed behavior of Heun and the working per-core grids for DPM-Solver++ and UniPC in the locked sweep.

A.2 Adams-Bashforth solvers violate the pointwise assumption

For a q -step Adams-Bashforth method, the update is an integral of an extrapolated integrand. Write the ODE in noise coordinate as

$$dx/dr = g(r, x(r)).$$

A q -step AB update has the form

$$x_{i+1} = x_i + h_i \sum_{\ell=0}^{q-1} b_{i,\ell} g_{i-\ell},$$

where $b_{i,\ell}$ depends on the nonuniform grid $(r_{i-q+1}, \dots, r_i, r_{i+1})$.

The local residual is not a pointwise derivative coefficient. It is the interpolation-kernel residual

$$e_i^{AB} = \int_{r_i}^{r_{i+1}} [g(r, x(r)) - I_i^{q-1} g(r)] dr,$$

$I_i^{q-1} g$ = extrapolation polynomial through $(r_i, g_i), (r_{i-1}, g_{i-1}), \dots, (r_{i-q+1}, g_{i-q+1})$.

Using interpolation error,

$$e_i^{AB} = (1/q!) \int_{r_i}^{r_{i+1}} g^{(q)}(xi_i) \prod_{\ell=0}^{q-1} (r - r_{i-\ell}) dr.$$

The product contains previous nodes. Therefore the residual depends on history:

$$e_i^{AB} = e_i^{AB}(r_{i+1}, r_i, r_{i-1}, \dots, r_{i-q+1}; g).$$

For nonuniform AB-2, with current step h_i and previous step h_{i-1} ,

$$\begin{aligned} e_i^{AB2} &= [h_i^3/6 + h_{i-1} h_i^2/4] g''(xi_i) \\ &= [h_i^2(2 h_i + 3 h_{i-1})/12] g''(xi_i). \end{aligned}$$

Thus DEIS tAB-2 cannot be certified by a scalar pointwise $d_{DEIS}(r)$ unless the previous spacing is fixed. If the grid changes, d_{DEIS} changes with it.

A.3 Corrected multistep certificate

For a q -step AB solver, replace the separable discretization term by a coupled sequence functional. Let the grid be

$$\text{Pi}_N = \{r_0 > r_1 > \dots > r_N\} \quad \text{or} \quad \{\sigma_0 > \sigma_1 > \dots > \sigma_N\}.$$

Define the amplified AB residual density empirically or analytically as

$$D_s^{\text{AB}}(\text{Pi}_N) = \sum_{i=q-1}^{N-1} A_i^2 \left\| K_{\{s,i\}}^q(g; r_{\{i+1\}}, \dots, r_{\{i-q+1\}}) \right\|^2.$$

Here A_i is terminal amplification and $K_{\{s,i\}}^q$ is the solver-specific extrapolation kernel residual. For DEIS tAB-2, K depends on $(h_i, h_{\{i-1\}})$ and on the integrand history $(g_i, g_{\{i-1\}})$; for PNDM/PLMS-4 it depends on four previous integrand evaluations and three or four previous spacings.

The corrected KL/NLL certificate is

$$\begin{aligned} Q_s^{\text{AB}}(\rho, \text{Pi}_N) \\ = Q_0 + (c_{\text{stat}}/n) \int_0^B a(r)/\rho(r) \, dr \\ + c_{\text{disc}} D_s^{\text{AB}}(\text{Pi}_N). \end{aligned}$$

The schedule-density optimizer for the statistical term remains

$$\rho^*(r) \text{ proportional to } \sqrt{a(r)}.$$

The sampling grid optimizer is no longer pointwise. It is

$$\text{Pi}_{\{N,s\}}^* = \underset{\{\text{Pi}_N\}}{\text{argmin}} D_s^{\text{AB}}(\text{Pi}_N),$$

subject to monotone grid constraints and any $\sigma_{\text{min}}/\sigma_{\text{max}}$ bounds.

Equivalently, the accuracy-constrained form is

$$\begin{aligned} &\text{minimize } N \\ &\text{subject to } Q_s^{\text{AB}}(\rho, \text{Pi}_N) \leq \epsilon. \end{aligned}$$

This is the correct replacement for $m_s^*(r)$ proportional to $d_s(r)^{1/(p+1)}$ in AB-class solvers.

A.4 What to do with d_s in practice

- (a) History-dependent d_s is the correct primary form: use $d_s(r_i, \dots, r_{\{i-q\}})$ or $D_s^{\text{AB}}(\text{Pi}_N)$, then optimize the whole sequence.
- (b) Residualized d_s is a secondary diagnostic: subtract the extrapolation-kernel residual only if the remaining defect is pointwise. It is not the main DEIS/PLMS certificate.
- (c) The corrected form is a sequence-level program: use a coupled residual functional; the pointwise m^* formula is recovered only when the residual separates.

A.5 Updated Pareto-domination theorem

The original Pareto theorem is split by solver family.

Theorem A: single-step or pointwise-defect solvers.

Let the solver certificate be separable:

$$Q_s(\rho, m) = Q_0 + (c_{\text{stat}}/n) \int a(r)/\rho(r) dr + c_{\text{disc}} \int d_s(r)/m(r)^p dr.$$

Then

$$\rho^*(r) \text{ proportional to } \sqrt{a(r)}, \\ m_s^*(r) \text{ proportional to } d_s(r)^{1/(p+1)}$$

minimizes the certificate among all schedules and grids with the same resource budget. Hence, for every baseline grid in this single-step admissible class, the optimized scheduler weakly Pareto-dominates it under the same certified accuracy-efficiency criterion.

Theorem B: q-step Adams-Bashforth-class solvers.

Let the solver certificate be

$$Q_s^{\text{AB}}(\rho, \Pi_N) = Q_0 + (c_{\text{stat}}/n) \int a(r)/\rho(r) dr + c_{\text{disc}} D_s^{\text{AB}}(\Pi_N).$$

Then the Pareto optimizer is

$$\rho^*(r) \text{ proportional to } \sqrt{a(r)}, \\ \Pi_{\{N, s\}}^* = \operatorname{argmin}_{\{\Pi_N\}} D_s^{\text{AB}}(\Pi_N).$$

For every baseline AB grid Π_b with N_b steps, choosing $\epsilon = Q_s^{\text{AB}}(\rho_b, \Pi_b)$ yields an optimized sequence Π^* with

$$Q_s^{\text{AB}}(\rho^*, \Pi^*) \leq Q_s^{\text{AB}}(\rho_b, \Pi_b), \\ N(\Pi^*) \leq N_b.$$

Thus Pareto domination still holds, but only for the corrected sequence-level optimizer. The pointwise grid $m_s^*(r)$ proportional to $d_s(r)^{1/(p+1)}$ is not a theorem for AB-class solvers.

A.6 Empirical implication

The DEIS result is not an anomaly. It is the expected failure mode of applying a single-step certificate to a multistep extrapolation method. A per-core $d_{\text{DEIS}}(r)$ estimated from one-step DEIS-vs-reference error measures an extrapolation-kernel residual under the calibration grid. When the optimizer changes the grid, that residual changes. The scalar density $d_{\text{DEIS}}(r)$ is therefore not an invariant local stiffness profile.

The locked sweep should be read as follows:

- Heun, DPM-Solver++, and UniPC pass the pointwise or effectively pointwise certificate test.
- DEIS tAB-2 fails because its leading defect is history-coupled.
- PNDM/PLMS-4 is predicted to require the same AB-class correction, although it was not directly tested in the locked per-core sweep.

A shared d_{Heun} grid can improve DEIS because it reintroduces low-sigma sensitivity, but it is not a DEIS per-core certificate. The corrected DEIS certificate must optimize the AB residual over the whole sequence.

A.7 Replacement statement for the main text

Replace the unqualified claim

$$m_s^*(r) \text{ proportional to } d_s(r)^{1/(p+1)} \text{ for each solver core } s$$

by

For pointwise-defect solvers, $m_s^*(r)$ proportional to $d_s(r)^{1/(p+1)}$.
For q-step Adams-Bashforth-class solvers, replace $d_s(r)$ by the
history-dependent residual $D_s^{AB}(P_i_N)$, and compute the optimal grid
by sequence optimization.

This correction preserves the accuracy-constrained efficiency principle and removes the false per-core DEIS
implication.