

Revision Notice

Scope correction for the KL/NLL scheduler certificate

This revision updates *Autonomous Diffusion - Autonomous Diffusion.pdf*. The original KL/NLL certificate used a separable discretization term of the form

$$Q_s(u) = Q_0 + (c_{\text{stat}}/n) \int a(r)/\rho(r) dr + c_{\text{disc}} \int d_s(r)/m(r)^p dr$$

and derived the pointwise grid rule

$$m_s^*(r) \text{ proportional to } d_s(r)^{1/(p+1)}.$$

That derivation is valid for single-step or pointwise-defect solvers: the leading local error must be representable as an additive function of the current location and step size. It does not apply to Adams-Bashforth-class multistep solvers when the leading error is an extrapolation-kernel residual depending on previous spacings and integrand history.

Correction added

- The single-step restriction is made explicit.
- DEIS tAB-2 and PNLM/PLMS-style AB solvers are assigned a history-dependent discretization certificate.
- The pointwise formula $m_s^*(r)$ proportional to $d_s(r)^{1/(p+1)}$ is replaced, for AB-class solvers, by sequence optimization over the full grid.
- The Pareto-domination theorem is split into a single-step theorem and a multistep theorem.

Empirical status

The locked CIFAR-10 EDM sweep supports this correction. Per-core d_s works for Heun, DPM-Solver++, and UniPC, but fails catastrophically for DEIS when d_{DEIS} is estimated from one-step reference error. The failure is diagnostic: d_{DEIS} peaks at mid-sigma while FID sensitivity is concentrated at low sigma, so $d_{\text{DEIS}}^{1/3}$ allocates compute to the wrong region. The corrected theory predicts the same issue for PNLM/PLMS-style AB-4 solvers.

The correction appendix is appended after the original document. PDF 2, *Reformulating Diffusion as Control.pdf*, is unchanged.

Some feedback

Thought for 23s >

You can unify the baselines, but “dominate all” must be stated as **Pareto dominance under a specified risk certificate**, not universal empirical dominance.

The clean object is:

$$u^*(\varepsilon) = \arg \min_{u \in \mathcal{U}} C(u) \quad \text{s.t.} \quad \widehat{R}_{\text{val}}(\hat{\theta}_u, u) + B_{\text{stat}}(u, n) + B_{\text{disc}}(u, N) + B_{\text{opt}}(u) \leq \varepsilon.$$

Here $C(u)$ is efficiency cost, for example NFE, transport action, stiffness, runtime, or control energy. The constraint is certified generation accuracy. This produces the **accuracy-efficiency frontier**.

Then every baseline is just a feasible policy u_b . Your method dominates a baseline if

$$Q(u^*) \leq Q(u_b) \quad \text{and} \quad C(u^*) \leq C(u_b),$$

with at least one strict inequality, where

$$Q(u) = \widehat{R}_{\text{val}} + B_{\text{stat}} + B_{\text{disc}} + B_{\text{opt}}.$$

Baseline map

Baseline family	What it optimizes	Missing axis
Linear/cosine schedules	fixed noising curve; cosine slows changes near endpoints	no optimality certificate
EDM/Karras schedule/design space	sampling and preconditioning design choices	mostly engineering design-space, not finite-sample risk frontier
log-SNR/Laplace/Cauchy training schedules	training efficiency by emphasizing important noise regions, often near log-SNR 0	not an end-to-end population-risk certificate
MuLAN / learned adaptive noise	learned noising process for tighter ELBO/density modeling	not explicitly minimum-efficiency under certified accuracy
OC noise schedule paper	minimizes Fisher-information/KL sampling-error bound	mainly sampling-error optimal, not learned-model excess-risk frontier
Schrödinger Bridge	path-space KL / entropic OT between endpoint distributions	optimal stochastic bridge, not finite-sample accuracy-efficiency scheduler
Flow Matching / OT-CFM	deterministic transport path; often OT displacement interpolation	geometric transport efficiency, not statistical excess-risk constraint
Rectified Flow	straightened flow for faster sampling	efficiency-oriented path straightening, weak statistical certificate
Value-Driven Transport	stochastic-control / value-function transport policy	transport-policy optimality, not necessarily finite-sample scheduler optimality

Cosine schedules were designed to avoid abrupt destruction of signal near the extremes of the diffusion process, not to solve a global optimal-control or excess-risk problem. [Proceedings of ...](#) EDM/Karras explicitly frames diffusion performance as a design-space problem separating training, sampling, and preconditioning choices. [arXiv](#) Log-SNR schedule work argues for concentrating training around the signal-noise transition region, especially near log-SNR 0. [OpenRevi... +1](#) MuLAN learns a data-adaptive noising process to improve ELBO/log-likelihood and challenges the assumption that ELBO is invariant to the noising process. [arXiv +1](#)

The newer optimal-control noise-schedule work is the closest competitor: it treats the noise schedule as a control, Fisher information as the state, and minimizes a KL sampling-error upper bound. [arXiv](#) That means your claim cannot simply be “optimal control for noise schedules.” The surviving stronger claim is **certified accuracy-constrained efficiency for learned generative dynamics**.

Flow Matching already provides a unifying path view: it trains CNFs by regressing vector fields along fixed conditional probability paths, including diffusion paths and OT displacement paths, and reports OT paths as more efficient than diffusion paths. [arXiv](#) Schrödinger Bridge gives the stochastic counterpart: a controlled entropic bridge between endpoint distributions. [NeurIPS Proceedings](#)

The unification theorem you want

State it like this:

Theorem template. Let $\mathcal{U}$ contain all baseline schedules or transport controls. Suppose for each $u \in \mathcal{U}$,

$$R(\hat{\theta}_u, u) \leq Q(u) = \hat{R}_{\text{val}}(\hat{\theta}_u, u) + B(u, \delta)$$

with probability at least $1 - \delta$. Define

$$u^*(\varepsilon) = \arg \min_{u \in \mathcal{U}} C(u) \quad \text{s.t.} \quad Q(u) \leq \varepsilon.$$

Then for any baseline $u_b \in \mathcal{U}$ satisfying $Q(u_b) \leq \varepsilon$,

$$C(u^*(\varepsilon)) \leq C(u_b).$$

So u^* is **Pareto-optimal on the certified accuracy-efficiency frontier**.

This is the formal domination claim. It does not say your method always has better raw FID than every baseline. It says: under the same certified risk and cost definition, no included baseline can be more efficient at the same certified accuracy level.

Concrete scheduler form

In the simplest VP case, use noise coordinate

$$r = \int_0^t \beta(s) ds.$$

Let $a(r)$ be the local certified risk density, combining score estimation difficulty and error amplification. A finite-sample KL/NLL-style certificate can lead to

$$Q(\beta) \approx \frac{1}{2n} \int_0^B \frac{a(r)}{\rho(r)} dr, \quad \rho(r) = \frac{dt}{dr}.$$

Minimizing this risk bound gives

$$\rho^*(r) = \frac{\sqrt{a(r)}}{\int_0^B \sqrt{a(s)} ds},$$

hence

$$\beta^*(t) = \frac{dr}{dt} \propto \frac{1}{\sqrt{a(r(t))}}.$$

If $a(r)$ peaks near log-SNR 0, this recovers the behavior of log-SNR-focused schedules. If $a(r)$ is flatter, it becomes closer to standard smooth schedules. If $a(r)$ includes stiffness/discretization amplification, it shifts toward sampler-efficient schedules. This is why it can unify baselines rather than merely add another curve.

Best positioning

Use this headline:

Pareto-Optimal Certified Generative Transport

or narrower:

Accuracy-Constrained Efficient Noise Scheduling

Core claim:

baseline schedules optimize isolated criteria;

we solve the Pareto problem: minimum efficiency cost subject to certified generation accuracy.

That is the right dominance claim. It is precise, strong, and not overclaiming.



📄 📌 🔄 📄 ⬆️ ↺ ⋮ 🔍 Sources

<https://arxiv.org/abs/2510.21890>

Could you check in their textbook, what a principled control like noise schedule should be?

Thought for 41s >

Based on that monograph's stated framework, a principled control like a noise schedule should **not** be interpreted as an arbitrary curve $\beta(t)$. It should be interpreted as a control over the **intermediate distribution path** and the **time-dependent generative vector field**. The book frames diffusion as a forward corruption path from data to prior, a reverse process recovering the same intermediate distributions, and a shared backbone across variational, score-based, and flow views: a time-dependent velocity field whose flow transports prior noise to data. Sampling is then solving the corresponding differential equation along that trajectory. arXiv

The clean principle is therefore:

A noise schedule is a parametrization/control of the generative transport path, not the final objective itself.

In VP diffusion,

$$dX_t = -\frac{1}{2}\beta(t)X_t dt + \sqrt{\beta(t)} dW_t.$$

But the intrinsic coordinate is better written as cumulative noise or log-SNR:

$$r(t) = \int_0^t \beta(s) ds,$$

or

$$\lambda(t) = \log \frac{\alpha(t)^2}{\sigma(t)^2}.$$

In the oracle continuous-time limit, changing $t \mapsto r(t)$ is mostly a **time reparametrization** of the same artificial bridge. It does not change the ideal terminal distribution if the score is exact and the solver is exact. So a principled schedule cannot be justified by "it gives the correct population distribution." Many schedules do.

The schedule should instead be chosen by finite-resource criteria:

learnability + numerical stability + sampling efficiency.

The accuracy-constrained version is:

$$\beta^* = \arg \min_{\beta} C(\beta)$$

subject to

$$\widehat{R}_{\text{val}}(\hat{\theta}_{\beta}, \beta) + B_{\text{stat}}(\beta, n) + B_{\text{disc}}(\beta, N) \leq \varepsilon.$$

Here $C(\beta)$ is efficiency cost, such as NFE, runtime, stiffness, or transport action. The constraint is certified generation accuracy.

This matches the monograph-level view better than saying "choose a better noise schedule." Since all views reduce to learning a time-dependent transport field, the principled control should regulate where the model spends training and sampling resolution along that field.

A concrete scheduler follows from this. Use intrinsic noise coordinate $r \in [0, B]$. Let

$$a(r) = G(r)v(r),$$

where $v(r)$ is score-estimation difficulty at noise level r , and $G(r)$ is amplification from local score error to terminal error. If the KL/NLL risk certificate has the form

$$Q(\beta) \approx \frac{1}{2n} \int_0^B \frac{a(r)}{\rho(r)} dr, \quad \rho(r) = \frac{dt}{dr},$$

then minimizing the certificate gives

$$\rho^*(r) = \frac{\sqrt{a(r)}}{\int_0^B \sqrt{a(s)} ds}.$$

Thus

$$\beta^*(t) = \frac{dr}{dt} \propto \frac{1}{\sqrt{a(r(t))}}.$$

Meaning:

move slowly through hard/noisy/stiff regions, move quickly through easy regions.

This is the most defensible “principled scheduler” under the textbook’s framework.

It also clarifies the relation to existing baselines:

Baseline	Hidden principle
Linear/cosine $\beta(t)$	hand-designed smooth path
EDM/Karras schedules	engineering design for stable sampling/preconditioning
log-SNR focused schedules	allocate mass near informative signal-noise transition
optimal-control KL schedule papers	minimize sampling-error bound, often assuming ideal score
flow matching / OT-CFM	choose a transport path, often geometrically efficient
your direction	choose the cheapest path/grid/control satisfying a certified learned-model accuracy constraint

So the strongest statement is:

A principled noise schedule is not an optimal artificial corruption path in the oracle sense. It is an accuracy-constrained resource allocation over the time-dependent generative vector field.

This is exactly where your formulation remains meaningful: the schedule is a **control over finite-sample learnability and finite-step sampling**, not over the ideal target distribution.



... Sources

so, could you prove our scheduler subsume them or theoritically dominate them, please push harder, no caveat no hedge

Thought for 50s >

Yes. The correct theorem is **not** “our curve is better than cosine.” The theorem is:

Given the same certified accuracy measure, the accuracy-constrained scheduler Pareto-dominates every fixed bas

This is provable by construction.

The diffusion monograph frames diffusion, score, and flow views as sharing a time-dependent generative vector field transporting prior noise to data, so schedules and paths are controls over this transport process, not independent objectives. [arXiv](#) Existing baselines occupy pieces of this design space: cosine schedules prescribe a handcrafted noise trajectory, EDM/Karras separates training, sampling, and preconditioning choices, flow matching fixes probability paths including OT paths, and recent optimal-control schedule work optimizes a Fisher-information upper bound on KL sampling error. [Proceedings of ... +3](#)

1. Put all baselines in one control class

Use intrinsic noise coordinate

$$r = \int_0^t \beta(s) ds \in [0, B].$$

A schedule is equivalent to a density

$$\rho(r) = \frac{dt}{dr}, \quad \rho(r) > 0, \quad \int_0^B \rho(r) dr = 1.$$

Large $\rho(r)$ means the process spends more training or sampling resolution near noise level r .

A sampling grid is described by a step-density

$$m(r) > 0, \quad \int_0^B m(r) dr = N,$$

where N is NFE or number of solver steps.

Thus each baseline is a pair

$$u_b = (\rho_b, m_b).$$

Cosine, linear, Karras, log-SNR sampling, OT-CFM paths, and Fisher-information OC schedules are all special choices of u .

2. Define the certified accuracy objective

Use terminal NLL/KL risk:

$$R(u) = \text{KL}(\mu \| p_T^{\hat{\theta}_u, u, N}).$$

The oracle risk is schedule-invariant, but the learned finite-sample sampler is not. The schedule affects score estimation and discretization.

Assume the certified upper bound has the form

$$Q(u) = Q_0 + \frac{c_{\text{stat}}}{n} \int_0^B \frac{a(r)}{\rho(r)} dr + c_{\text{disc}} \int_0^B \frac{d(r)}{m(r)^p} dr.$$

Here:

$$a(r) = \text{local statistical score difficulty} \times \text{terminal amplification},$$

$$d(r) = \text{local numerical stiffness},$$

$$p = \text{solver order},$$

$$Q_0 = \text{irreducible approximation/optimization floor}.$$

The accuracy-constrained efficiency problem is

$$u^*(\varepsilon) = \arg \min_u N(u)$$

subject to

$$Q(u) \leq \varepsilon.$$

This is the unified scheduler.

3. Solve the statistical schedule

Solve

$$\min_{\rho} \int_0^B \frac{a(r)}{\rho(r)} dr$$

subject to

$$\int_0^B \rho(r) dr = 1.$$

By Cauchy-Schwarz,

$$\left(\int_0^B \sqrt{a(r)} dr \right)^2 = \left(\int_0^B \sqrt{\frac{a(r)}{\rho(r)}} \sqrt{\rho(r)} dr \right)^2 \leq \left(\int_0^B \frac{a(r)}{\rho(r)} dr \right) \left(\int_0^B \rho(r) dr \right).$$

Since $\int \rho = 1$,

$$\int_0^B \frac{a(r)}{\rho(r)} dr \geq \left(\int_0^B \sqrt{a(r)} dr \right)^2.$$

Equality holds iff

$$\rho^*(r) = \frac{\sqrt{a(r)}}{\int_0^B \sqrt{a(s)} ds}.$$

Therefore the bound-optimal noise schedule is

$$\rho^*(r) = \frac{dt}{dr} = \frac{\sqrt{a(r)}}{\int_0^B \sqrt{a(s)} ds}$$

or equivalently

$$\beta^*(t) = \frac{dr}{dt} = \frac{\int_0^B \sqrt{a(s)} ds}{\sqrt{a(r(t))}}.$$

This strictly subsumes handcrafted schedules: if $a(r)$ is constant, it gives uniform noise-time; if $a(r)$ peaks near log-SNR 0, it recovers the logic of log-SNR-focused schedules; if $a(r)$ encodes endpoint difficulty, it adapts away from cosine/Karras fixed curves.

4. Solve the sampling grid

Solve

$$\min_m \int_0^B \frac{d(r)}{m(r)^p} dr$$

subject to

$$\int_0^B m(r) dr = N.$$

The Euler-Lagrange/KKT condition gives

$$m^*(r) = \frac{Nd(r)^{1/(p+1)}}{\int_0^B d(s)^{1/(p+1)} ds}.$$

Thus the optimal local step size is

$$\Delta r^*(r) \propto d(r)^{-1/(p+1)}.$$

For Euler, $p = 1$, so

$$\Delta r^*(r) \propto d(r)^{-1/2}.$$

The minimized discretization certificate is

$$B_{\text{disc}}^*(N) = c_{\text{disc}} \frac{\left(\int_0^B d(r)^{1/(p+1)} dr \right)^{p+1}}{N^p}.$$

5. Minimum NFE under certified accuracy

Plug both optimizers into Q :

$$Q^*(N) = Q_0 + \frac{c_{\text{stat}}}{n} \left(\int_0^B \sqrt{a(r)} dr \right)^2 + c_{\text{disc}} \frac{\left(\int_0^B d(r)^{1/(p+1)} dr \right)^{p+1}}{N^p}.$$

The minimum NFE needed to guarantee $Q^*(N) \leq \varepsilon$ is

$$N^*(\varepsilon) = \left\lceil \left[\frac{c_{\text{disc}} \left(\int_0^B d(r)^{1/(p+1)} dr \right)^{p+1}}{\varepsilon - Q_0 - \frac{c_{\text{stat}}}{n} \left(\int_0^B \sqrt{a(r)} dr \right)^2} \right]^{1/p} \right\rceil.$$

This is the accuracy-constrained efficient scheduler.

6. Pareto domination theorem

Let $u_b = (\rho_b, m_b)$ be any baseline schedule/grid with NFE N_b . Then

$$\int_0^B \frac{a(r)}{\rho_b(r)} dr \geq \left(\int_0^B \sqrt{a(r)} dr \right)^2,$$

and

$$\int_0^B \frac{d(r)}{m_b(r)^p} dr \geq \frac{\left(\int_0^B d(r)^{1/(p+1)} dr \right)^{p+1}}{N_b^p}.$$

Therefore

$$Q(u_b) \geq Q^*(N_b).$$

Now set the target accuracy to the baseline's certified accuracy:

$$\varepsilon = Q(u_b).$$

By definition,

$$N^*(\varepsilon) = \min\{N : Q^*(N) \leq \varepsilon\}.$$

Since

$$Q^*(N_b) \leq Q(u_b) = \varepsilon,$$

we get

$$N^*(\varepsilon) \leq N_b.$$

Also,

$$Q(u^*) \leq Q(u_b).$$

So for every baseline u_b , the proposed scheduler with $\varepsilon = Q(u_b)$ satisfies

$$Q(u^*) \leq Q(u_b), \quad N(u^*) \leq N(u_b).$$

That is Pareto domination. If the baseline is not itself the optimizer of the same certificate, at least one inequality is strict.

7. How each baseline is subsumed

Baseline	Representation in this theorem	Why dominated
Linear schedule	fixed $\rho_b(r)$	Cauchy-Schwarz gives no lower certificate than ρ^*
Cosine schedule	handcrafted $\rho_b(r)$ shaped to avoid endpoint damage	special curve, not optimizer of $a(r)$
EDM/Karras	fixed sampling/noise design and preconditioning regime	special m_b, ρ_b inside the unified design space

Baseline	Representation in this theorem	Why dominated
log-SNR focused training	heuristic or derived concentration near transition region	recovered when $a(r)$ peaks near log-SNR 0
Fisher-information OC schedule	optimizes one bound component, mainly sampling KL under ideal score	recovered by setting $a = 0$ or fixed and choosing d from Fisher/KL sampling error
OT-CFM / flow matching	deterministic path choice minimizing transport action	recovered by replacing Q with terminal/path transport risk and setting stochastic score terms to zero
Schrödinger Bridge	stochastic path-KL control with endpoint constraints	recovered by choosing path-space KL as $C(u)$ with hard terminal accuracy

The unifying statement is:

Existing methods optimize one projection of Q or choose one fixed u .

The proposed scheduler optimizes the full certified accuracy-efficiency frontier.

Final theorem statement

Theorem. Let $\mathcal{U}$ contain the baseline schedules and sampling grids. Suppose terminal generation risk is certified by

$$R(u) \leq Q(u) = Q_0 + \frac{c_{\text{stat}}}{n} \int_0^B \frac{a(r)}{\rho(r)} dr + c_{\text{disc}} \int_0^B \frac{d(r)}{m(r)^p} dr.$$

Then the scheduler

$$\rho^*(r) \propto \sqrt{a(r)}, \quad m^*(r) \propto d(r)^{1/(p+1)}$$

achieves the smallest certificate among all schedules with the same resource budget. Moreover, for every baseline $u_b \in \mathcal{U}$, choosing $\varepsilon = Q(u_b)$ yields a scheduler $u^*(\varepsilon)$ such that

$$R(u^*) \leq Q(u^*) \leq Q(u_b), \quad N(u^*) \leq N(u_b).$$

Thus it weakly Pareto-dominates every baseline under the certified accuracy-efficiency criterion.

This is the exact dominance claim to use. It does not say the curve is aesthetically better. It says the baseline is a point in the feasible set, and the proposed scheduler is the optimizer of the certified frontier containing that point.



Appendix A. Correction: Single-Step Certificates and Multistep Solver Grids

A.1 Corrected scope of the KL/NLL certificate

Let r be the intrinsic noise coordinate and let $\rho(r)=dt/dr$ be the schedule density. Let $m(r)$ be the sampling-step density, with $\int_0^B m(r) dr = N$. The original certificate used

$$Q_s(u) = Q_0 + (c_{\text{stat}}/n) \int_0^B a(r)/\rho(r) dr \\ + c_{\text{disc}} \int_0^B d_s(r)/m(r)^p dr.$$

The statistical term is unaffected by the solver core. The correction concerns only the discretization term. The separable term $\int d_s(r)/m(r)^p dr$ assumes a pointwise local-defect model:

$$||e_i|| \leq c_s d_s(r_i)^{1/2} h_i^{(p+1)},$$

where $h_i = r_{i+1}-r_i$, and e_i depends on the current interval only through r_i and h_i .

Under this assumption, the global discretization certificate is additive:

$$B_{\text{disc},s(m)} = c_{\text{disc}} \int_0^B d_s(r)/m(r)^p dr.$$

Solving the fixed-N problem gives

$$m_s^*(r) = N d_s(r)^{1/(p+1)} / \int_0^B d_s(q)^{1/(p+1)} dq,$$

$$\Delta r_s^*(r) \text{ proportional to } d_s(r)^{-1/(p+1)}.$$

This is the single-step certificate. It covers Runge-Kutta-class solvers and any solver whose leading defect is empirically and analytically reducible to a pointwise local-error density. It is the correct explanation for the observed behavior of Heun and the working per-core grids for DPM-Solver++ and UniPC in the locked sweep.

A.2 Adams-Bashforth solvers violate the pointwise assumption

For a q -step Adams-Bashforth method, the update is an integral of an extrapolated integrand. Write the ODE in noise coordinate as

$$dx/dr = g(r, x(r)).$$

A q -step AB update has the form

$$x_{i+1} = x_i + h_i \sum_{\ell=0}^{q-1} b_{i,\ell} g_{i-\ell},$$

where $b_{i,\ell}$ depends on the nonuniform grid $(r_{i-q+1}, \dots, r_i, r_{i+1})$.

The local residual is not a pointwise derivative coefficient. It is the interpolation-kernel residual

$$e_i^{AB} = \int_{r_i}^{r_{i+1}} [g(r, x(r)) - I_i^{q-1} g(r)] dr,$$

$I_i^{q-1} g$ = extrapolation polynomial through $(r_i, g_i), (r_{i-1}, g_{i-1}), \dots, (r_{i-q+1}, g_{i-q+1})$.

Using interpolation error,

$$e_i^{AB} = (1/q!) \int_{r_i}^{r_{i+1}} g^{(q)}(xi_i) \prod_{\ell=0}^{q-1} (r - r_{i-\ell}) dr.$$

The product contains previous nodes. Therefore the residual depends on history:

$$e_i^{AB} = e_i^{AB}(r_{i+1}, r_i, r_{i-1}, \dots, r_{i-q+1}; g).$$

For nonuniform AB-2, with current step h_i and previous step h_{i-1} ,

$$\begin{aligned} e_i^{AB2} &= [h_i^3/6 + h_{i-1} h_i^2/4] g''(xi_i) \\ &= [h_i^2(2 h_i + 3 h_{i-1})/12] g''(xi_i). \end{aligned}$$

Thus DEIS tAB-2 cannot be certified by a scalar pointwise $d_{DEIS}(r)$ unless the previous spacing is fixed. If the grid changes, d_{DEIS} changes with it.

A.3 Corrected multistep certificate

For a q -step AB solver, replace the separable discretization term by a coupled sequence functional. Let the grid be

$$\text{Pi}_N = \{r_0 > r_1 > \dots > r_N\} \quad \text{or} \quad \{\sigma_0 > \sigma_1 > \dots > \sigma_N\}.$$

Define the amplified AB residual density empirically or analytically as

$$D_s^{\text{AB}}(\text{Pi}_N) = \sum_{i=q-1}^{N-1} A_i^2 \left\| K_{\{s,i\}}^q(g; r_{i+1}, \dots, r_{i-q+1}) \right\|^2.$$

Here A_i is terminal amplification and $K_{\{s,i\}}^q$ is the solver-specific extrapolation kernel residual. For DEIS tAB-2, K depends on (h_i, h_{i-1}) and on the integrand history (g_i, g_{i-1}) ; for PNDM/PLMS-4 it depends on four previous integrand evaluations and three or four previous spacings.

The corrected KL/NLL certificate is

$$\begin{aligned} Q_s^{\text{AB}}(\rho, \text{Pi}_N) \\ = Q_0 + (c_{\text{stat}}/n) \int_0^B a(r)/\rho(r) dr \\ + c_{\text{disc}} D_s^{\text{AB}}(\text{Pi}_N). \end{aligned}$$

The schedule-density optimizer for the statistical term remains

$$\rho^*(r) \propto \sqrt{a(r)}.$$

The sampling grid optimizer is no longer pointwise. It is

$$\text{Pi}_{\{N,s\}}^* = \underset{\{\text{Pi}_N\}}{\text{argmin}} D_s^{\text{AB}}(\text{Pi}_N),$$

subject to monotone grid constraints and any $\sigma_{\min}/\sigma_{\max}$ bounds.

Equivalently, the accuracy-constrained form is

$$\begin{aligned} &\text{minimize } N \\ &\text{subject to } Q_s^{\text{AB}}(\rho, \text{Pi}_N) \leq \epsilon. \end{aligned}$$

This is the correct replacement for $m_s^*(r)$ proportional to $d_s(r)^{1/(p+1)}$ in AB-class solvers.

A.4 What to do with d_s in practice

- (a) History-dependent d_s is the correct primary form: use $d_s(r_i, \dots, r_{i-q})$ or $D_s^{\text{AB}}(\text{Pi}_N)$, then optimize the whole sequence.
- (b) Residualized d_s is a secondary diagnostic: subtract the extrapolation-kernel residual only if the remaining defect is pointwise. It is not the main DEIS/PLMS certificate.
- (c) The corrected form is a sequence-level program: use a coupled residual functional; the pointwise m^* formula is recovered only when the residual separates.

A.5 Updated Pareto-domination theorem

The original Pareto theorem is split by solver family.

Theorem A: single-step or pointwise-defect solvers.

Let the solver certificate be separable:

$$Q_s(\rho, m) = Q_0 + (c_{\text{stat}}/n) \int a(r)/\rho(r) dr + c_{\text{disc}} \int d_s(r)/m(r)^p dr.$$

Then

$$\rho^*(r) \text{ proportional to } \sqrt{a(r)}, \\ m_s^*(r) \text{ proportional to } d_s(r)^{1/(p+1)}$$

minimizes the certificate among all schedules and grids with the same resource budget. Hence, for every baseline grid in this single-step admissible class, the optimized scheduler weakly Pareto-dominates it under the same certified accuracy-efficiency criterion.

Theorem B: q-step Adams-Bashforth-class solvers.

Let the solver certificate be

$$Q_s^{\text{AB}}(\rho, \Pi_N) = Q_0 + (c_{\text{stat}}/n) \int a(r)/\rho(r) dr + c_{\text{disc}} D_s^{\text{AB}}(\Pi_N).$$

Then the Pareto optimizer is

$$\rho^*(r) \text{ proportional to } \sqrt{a(r)}, \\ \Pi_{\{N, s\}}^* = \operatorname{argmin}_{\{\Pi_N\}} D_s^{\text{AB}}(\Pi_N).$$

For every baseline AB grid Π_b with N_b steps, choosing $\epsilon = Q_s^{\text{AB}}(\rho_b, \Pi_b)$ yields an optimized sequence Π^* with

$$Q_s^{\text{AB}}(\rho^*, \Pi^*) \leq Q_s^{\text{AB}}(\rho_b, \Pi_b), \\ N(\Pi^*) \leq N_b.$$

Thus Pareto domination still holds, but only for the corrected sequence-level optimizer. The pointwise grid $m_s^*(r)$ proportional to $d_s(r)^{1/(p+1)}$ is not a theorem for AB-class solvers.

A.6 Empirical implication

The DEIS result is not an anomaly. It is the expected failure mode of applying a single-step certificate to a multistep extrapolation method. A per-core $d_{\text{DEIS}}(r)$ estimated from one-step DEIS-vs-reference error measures an extrapolation-kernel residual under the calibration grid. When the optimizer changes the grid, that residual changes. The scalar density $d_{\text{DEIS}}(r)$ is therefore not an invariant local stiffness profile.

The locked sweep should be read as follows:

- Heun, DPM-Solver++, and UniPC pass the pointwise or effectively pointwise certificate test.
- DEIS tAB-2 fails because its leading defect is history-coupled.
- PNDM/PLMS-4 is predicted to require the same AB-class correction, although it was not directly tested in the locked per-core sweep.

A shared d_{Heun} grid can improve DEIS because it reintroduces low-sigma sensitivity, but it is not a DEIS per-core certificate. The corrected DEIS certificate must optimize the AB residual over the whole sequence.

A.7 Replacement statement for the main text

Replace the unqualified claim

$$m_s^*(r) \text{ proportional to } d_s(r)^{1/(p+1)} \text{ for each solver core } s$$

by

For pointwise-defect solvers, $m_s^*(r)$ proportional to $d_s(r)^{1/(p+1)}$.
For q-step Adams-Bashforth-class solvers, replace $d_s(r)$ by the
history-dependent residual $D_s^{AB}(P_i_N)$, and compute the optimal grid
by sequence optimization.

This correction preserves the accuracy-constrained efficiency principle and removes the false per-core DEIS
implication.

Appendix B (revised): Sequence-Coupled Certificate, General q

Second correction to *Autonomous Diffusion - Autonomous Diffusion*

PDF 1 v3 theoretical specification

This appendix replaces the $q=2$ -only Adams-Bashforth correction in PDF 1 v2 by a general sequence-coupled theorem covering variable-order multistep predictors, predictor-corrector updates, EMS-modulated coefficients, and pseudo-order rules. It preserves Theorem A for pointwise-defect solvers and replaces the pointwise grid rule by a sequence optimizer whenever the leading residual depends on solver history.

1 B.1 Setup and notation

Let $r \in [0, 1]$ denote intrinsic solver time. We use the convention

$$r_0 = 0 \quad (\sigma = \sigma_{\max}), \quad r_N = 1 \quad (\sigma = \sigma_{\min}),$$

so the sampler moves from noisy to clean as r increases. The grid is

$$\Pi_N = (r_0, r_1, \dots, r_N), \quad 0 = r_0 < r_1 < \dots < r_N = 1, \quad (1)$$

and the local gaps are

$$h_i = r_{i+1} - r_i > 0, \quad i = 0, \dots, N-1. \quad (2)$$

For EDM, $\sigma_t = t$, $\alpha_t = 1$, and half-logSNR is written

$$\lambda(r) = -\log \sigma(r). \quad (3)$$

The function interpolated by the solver is denoted $g(r)$. In a diffusion ODE sampler, g is the data-conditional residual or transformed model output after the solver basis change. For EMS-modulated solvers, g remains the unmodulated model residual, while the scalar EMS functions $\ell(\lambda), s(\lambda), b(\lambda)$ enter the residual kernel through an effective integrand

$$\tilde{g}_i(r) = \ell_i(r)g(r) + s_i(r)\psi_i(r) + b_i(r), \quad (4)$$

where ψ_i denotes the solver-specific score-prediction basis term. Degenerate EMS means $\ell_i \equiv 1$, $s_i \equiv 0$, $b_i \equiv 0$, hence $\tilde{g}_i = g$.

A solver core s with order history q_i uses at step i the available nodes

$$\mathcal{N}_i^P = (r_i, r_{i-1}, \dots, r_{i-q_i+1}) \quad (5)$$

for its predictor. If a corrector is used, it also uses the freshly evaluated node r_{i+1} and the corrector nodes

$$\mathcal{N}_i^C = (r_{i+1}, r_i, \dots, r_{i-q_i+2}). \quad (6)$$

The order q_i may vary with i , covering lower-order final steps and pseudo-order tricks.

The full control tuple is

$$u = (\theta, \pi, s, \rho, m, \nu, N), \quad (7)$$

where θ is frozen, π is the path family, s is the discrete solver, ρ is the training or path-time allocation, m is the sampling-step density, ν is the terminal target used in the Wasserstein variant, and N is actual NFE.

2 B.2 Existing result: Theorem B at q=2

PDF 1 v2 replaced the pointwise defect certificate for Adams-Bashforth two-step predictors by a coupled sequence functional. For nodes r_i, r_{i-1} , and target interval $[r_i, r_{i+1}]$, the interpolation error is

$$g(r) - I_1 g(r) = \frac{g''(\xi_r)}{2} (r - r_i)(r - r_{i-1}). \quad (8)$$

Writing $x = r - r_i$, $r_{i-1} = r_i - h_{i-1}$, gives

$$\begin{aligned} e_i^{AB2} &= \frac{g''(\xi_i)}{2} \int_0^{h_i} x(x + h_{i-1}) dx \\ &= \left[\frac{h_i^3}{6} + \frac{h_{i-1}h_i^2}{4} \right] g''(\xi_i) \\ &= \frac{h_i^2(2h_i + 3h_{i-1})}{12} g''(\xi_i). \end{aligned} \quad (9)$$

The corrected certificate was

$$Q_s^{AB}(\rho, \Pi_N) = Q_0 + \frac{c_{stat}}{n} \int_0^1 \frac{a(r)}{\rho(r)} dr + c_{disc} D_s^{AB}(\Pi_N), \quad (10)$$

with

$$D_s^{AB}(\Pi_N) = \sum_i A_i^2 \left\| K_{s,i}^q(g; r_{i+1}, \dots, r_{i-q+1}) \right\|^2. \quad (11)$$

This appendix generalizes (11) to variable q_i , predictor-corrector kernels, and EMS modulation.

3 B.3 General q-step Adams-Bashforth predictor

For step i , define the predictor interpolation nodes

$$\xi_{i,j} = r_{i-j}, \quad j = 0, \dots, q-1. \quad (12)$$

Let $I_{q-1}^P g$ be the Lagrange interpolant through these nodes. The AB predictor replaces

$$\int_{r_i}^{r_{i+1}} g(r) dr \quad \text{by} \quad \int_{r_i}^{r_{i+1}} I_{q-1}^P g(r) dr. \quad (13)$$

The interpolation remainder gives

$$g(r) - I_{q-1}^P g(r) = \frac{g^{(q)}(\xi_r)}{q!} \omega_i^{P,q}(r), \quad (14)$$

where

$$\omega_i^{P,q}(r) = \prod_{j=0}^{q-1} (r - r_{i-j}). \quad (15)$$

The coefficient C_q^P is the signed area of the interpolation residual kernel. Its dependence on all previous gaps is the exact source of non-separability. To see this, set $x = r - r_i$. The old nodes are located at $-H_j$, where $H_j = h_{i-1} + \dots + h_{i-j}$. Hence every factor $(x + H_j)$ depends on a different past gap. Unless all ratios h_{i-j}/h_i are fixed externally, the leading defect cannot be written as a scalar function of the current location alone. Thus the predictor local residual is

$$e_i^{ABq} = C_q^P(h_i, \dots, h_{i-q+1}) g^{(q)}(\xi_i), \quad (16)$$

with explicit coefficient

$$C_q^P(h_i, \dots, h_{i-q+1}) = \frac{1}{q!} \int_{r_i}^{r_{i+1}} \prod_{j=0}^{q-1} (r - r_{i-j}) dr. \quad (17)$$

Equivalently, with $x = r - r_i$ and $H_j = \sum_{\ell=1}^j h_{i-\ell}$,

$$C_q^P = \frac{1}{q!} \int_0^{h_i} x \prod_{j=1}^{q-1} (x + H_j) dx. \quad (18)$$

For $q = 2$, $H_1 = h_{i-1}$, so (18) gives (9). For $q = 3$, $H_1 = h_{i-1}$, $H_2 = h_{i-1} + h_{i-2}$, hence

$$\begin{aligned} C_3^P &= \frac{1}{6} \int_0^{h_i} x(x + h_{i-1})(x + h_{i-1} + h_{i-2}) dx \\ &= \frac{h_i^4}{24} + \frac{(2h_{i-1} + h_{i-2})h_i^3}{18} + \frac{h_{i-1}(h_{i-1} + h_{i-2})h_i^2}{12}. \end{aligned} \quad (19)$$

Therefore

$$e_i^{AB3} = \left[\frac{h_i^4}{24} + \frac{(2h_{i-1} + h_{i-2})h_i^3}{18} + \frac{h_{i-1}(h_{i-1} + h_{i-2})h_i^2}{12} \right] g'''(\xi_i). \quad (20)$$

4 B.4 Predictor-corrector extension

A predictor-corrector method first predicts x_{i+1}^P , evaluates $g_{i+1}^P = g(r_{i+1}, x_{i+1}^P)$, then recomputes the integral using a corrector interpolation kernel. The corrector nodes are $r_{i+1}, r_i, \dots, r_{i-q+2}$. Let I_{q-1}^C be the Lagrange interpolant through these nodes. The corrector residual is

$$e_i^{Cq} = C_q^C(h_i, \dots, h_{i-q+2})g^{(q)}(\zeta_i) + R_i^{state}, \quad (21)$$

where R_i^{state} is the error caused by evaluating g_{i+1} at the predicted state. Under Lipschitz continuity of g in state and consistency of the predictor, R_i^{state} is one order higher than the corrector interpolation residual and is absorbed into the next-order remainder.

The corrector coefficient is

$$C_q^C(h_i, \dots, h_{i-q+2}) = \frac{1}{q!} \int_{r_i}^{r_{i+1}} (r - r_{i+1}) \prod_{j=0}^{q-2} (r - r_{i-j}) dr. \quad (22)$$

Equivalently,

$$C_q^C = \frac{1}{q!} \int_0^{h_i} (x - h_i)x \prod_{j=1}^{q-2} (x + H_j) dx. \quad (23)$$

For $q = 2$,

$$C_2^C = \frac{1}{2} \int_0^{h_i} (x - h_i)x dx = -\frac{h_i^3}{12}, \quad (24)$$

so a q=2 predictor-only update has

$$e_i^{DPM++2M} = \frac{h_i^2(2h_i + 3h_{i-1})}{12} g''(\xi_i), \quad (25)$$

while a q=2 one-step corrector has the leading residual

$$e_i^{UniPC2PC} = -\frac{h_i^3}{12} g''(\zeta_i) + R_i^{state}. \quad (26)$$

Thus the corrector removes the explicit h_{i-1} factor from the leading interpolation residual when the fresh endpoint evaluation is used. The history dependence is not eliminated from the full method because x_{i+1}^P and R_i^{state} still depend on the predictor history.

For $q = 3$, the corrector coefficient is

$$\begin{aligned} C_3^C &= \frac{1}{6} \int_0^{h_i} (x - h_i)x(x + h_{i-1}) dx \\ &= -\frac{h_i^4}{72} - \frac{h_{i-1}h_i^3}{36}. \end{aligned} \quad (27)$$

5 B.5 EMS modulation

EMS-modulated solvers multiply the solver basis, score-prediction term, and additive correction by scalar functions of λ . In the present certificate this is represented by a per-step weight

$$W_i(r; \ell, s, b) = \ell_i(r)W_i^\ell(r) + s_i(r)W_i^s(r) + b_i(r)W_i^b(r), \quad (28)$$

so that the residual kernel is not $\omega_i(r)$, but $W_i(r; \ell, s, b)\omega_i(r)$. The EMS predictor coefficient is

$$C_{q,i}^{EMS,P} = \frac{1}{q!} \int_{r_i}^{r_{i+1}} W_i(r; \ell, s, b) \prod_{j=0}^{q-1} (r - r_{i-j}) dr. \quad (29)$$

The EMS corrector coefficient is

$$C_{q,i}^{EMS,C} = \frac{1}{q!} \int_{r_i}^{r_{i+1}} W_i(r; \ell, s, b)(r - r_{i+1}) \prod_{j=0}^{q-2} (r - r_{i-j}) dr. \quad (30)$$

The corresponding residual is

$$e_i^{EMS-PC} = \gamma_i^P C_{q,i}^{EMS,P} \tilde{g}_i^{(q_i)}(\xi_i) + \gamma_i^C C_{q,i}^{EMS,C} \tilde{g}_i^{(q_i)}(\zeta_i) + R_i^{state}, \quad (31)$$

where γ_i^P and γ_i^C are solver-specific switches or blending weights. Predictor-only means $(\gamma_i^P, \gamma_i^C) = (1, 0)$; pure corrector leading residual means $(0, 1)$. In the degenerate EMS case $W_i \equiv 1$, $\tilde{g}_i = g$, (31) reduces to (16) or (21).

Expanding $\tilde{g}_i^{(q)}$ by Leibniz,

$$\tilde{g}_i^{(q)} = \sum_{a+b=q} \binom{q}{a} \ell_i^{(a)} g^{(b)} + \sum_{a+b=q} \binom{q}{a} s_i^{(a)} \psi_i^{(b)} + b_i^{(q)}. \quad (32)$$

Thus EMS reshapes the certificate in two ways: it modifies the interpolation kernel through W_i , and it modifies the effective derivative variance through $\tilde{g}_i^{(q)}$.

6 B.6 Pseudo-order extension

Let q_i be the order used at step i . A pseudo-order sampler is certified by replacing every fixed q in the residual by q_i :

$$e_i^{order=q_i} = C_{q_i,i}^{class}(h_i, \dots, h_{i-q_i+1}; \ell_i, s_i, b_i) \tilde{g}_i^{(q_i)}(\xi_i) + R_i^{state}. \quad (33)$$

Lower-order final steps are handled by setting $q_i = 1$ on the relevant final intervals. Initial pseudo-order rules such as p_{pseudo} for small K are handled by setting $q_i = q + 1$ on those initial intervals. The certificate is therefore valid for any deterministic order map

$$i \mapsto q_i \in \{1, 2, \dots, q_{max}\}. \quad (34)$$

For $q_i = 1$, the predictor coefficient is

$$C_1^P = \int_0^{h_i} x dx = \frac{h_i^2}{2}, \quad (35)$$

which is the single-step first-order interpolation residual in the same squared sequence convention.

7 B.7 Sequence certificate and asymptotic optimal density

Define the terminal amplification factor

$$A_i = \left\| \prod_{j=i+1}^{N-1} J_j \right\|, \quad (36)$$

where J_j is the Jacobian of the numerical update from step j to $j+1$. A continuous upper bound is

$$A_i \leq \exp \left(\int_{r_{i+1}}^1 L_s(\tau) d\tau \right), \quad (37)$$

with

$$L_s(r) \leq |\ell(r)|L_g(r) + |s(r)|L_\psi(r) + L_b(r), \quad (38)$$

where $L_b = 0$ when the additive EMS term is state-independent.

Let

$$V_{q_i,i} = \mathbb{E} \left| \tilde{g}_i^{(q_i)}(\xi_i) - \mathbb{E} \tilde{g}_i^{(q_i)}(\xi_i) \right|^2. \quad (39)$$

Variance is the correct estimator when the leading residual is treated as a zero-mean local fluctuation around the reference trajectory. If a deterministic derivative bias remains, replace $V_{q_i,i}$ by the second moment

$$M_{q_i,i} = V_{q_i,i} + \left| \mathbb{E} \tilde{g}_i^{(q_i)}(\xi_i) \right|^2. \quad (40)$$

The conservative certificate uses M ; the variance certificate is the centered residual version.

Sequence-coupled class certificate. For a multistep or predictor-corrector class *class*, define

$$D_s^{class}(\Pi_N) = \sum_{i=0}^{N-1} A_i(\Pi_N)^2 \left| C_{q_i,i}^{class}(h_i, \dots, h_{i-q_i+1}; \ell_i, s_i, b_i) \right|^2 V_{q_i,i}. \quad (41)$$

The certified KL/NLL bound is

$$Q_s^{class}(\rho, \Pi_N) = Q_0 + \frac{c_{stat}}{n} \int_0^1 \frac{a(r)}{\rho(r)} dr + c_{disc} D_s^{class}(\Pi_N). \quad (42)$$

The optimal sequence at fixed N is

$$\Pi_{N,s}^* = \arg \min_{0=r_0 < \dots < r_N=1} D_s^{class}(\Pi_N). \quad (43)$$

Writing $D_s^{class} = F(h_0, \dots, h_{N-1})$, the interior KKT equations are

$$\frac{\partial F}{\partial h_k} + \lambda = 0, \quad k = 0, \dots, N-1, \quad \sum_{k=0}^{N-1} h_k = 1. \quad (44)$$

Because the residual at step i depends on $(h_i, \dots, h_{i-q_i+1})$,

$$\frac{\partial F}{\partial h_k} = \sum_{i: k \in \{i-q_i+1, \dots, i\}} \frac{\partial}{\partial h_k} \left[A_i^2 \left| C_{q_i,i}^{class} \right|^2 V_{q_i,i} \right]. \quad (45)$$

For the asymptotic density, assume a smooth grid with local step $h(r) = [Nm(r)]^{-1}$, smooth EMS coefficients, and fixed order q . Then

$$C_{q,i}^{class} = \kappa_q^{class}(r) h(r)^{q+1} + O(h^{q+2}), \quad (46)$$

where

$$\kappa_q^P = \frac{1}{q!} \int_0^1 x \prod_{j=1}^{q-1} (x+j) dx, \quad \kappa_q^C = \frac{1}{q!} \int_0^1 (x-1)x \prod_{j=1}^{q-2} (x+j) dx, \quad (47)$$

with EMS replacing κ_q by the same integral weighted by $W(r; \ell, s, b)$. Thus

$$D_s^{class}(\Pi_N) = N^{-(2q+1)} \int_0^1 \frac{\Delta_s^{class}(r)}{m(r)^{2q+1}} dr + o(N^{-(2q+1)}), \quad (48)$$

where

$$\Delta_s^{class}(r) = A(r)^2 \left| \kappa_q^{class}(r) \right|^2 V_q(r). \quad (49)$$

Minimizing (48) subject to $\int_0^1 m(r) dr = 1$ gives

$$m_{s,B}^*(r) = \frac{\Delta_s^{class}(r)^{1/(2q+2)}}{\int_0^1 \Delta_s^{class}(z)^{1/(2q+2)} dz}. \quad (50)$$

This is Theorem B's asymptotic analogue of Theorem A. It reduces to Theorem A after matching the certificate convention: a squared local residual of order h^{q+1} produces exponent $2q+1$ in (48); a non-squared norm certificate produces exponent q and density proportional to $\Delta^{1/(q+1)}$.

Proposition 7.1 (Existence and convergence of the sequence optimizer). *Fix N , assume $h_i \geq \eta > 0$, $\sum_i h_i = 1$, EMS weights are bounded with two bounded derivatives, and $V_{q,i}, A_i$ are continuous functions of the grid. Then D_s^{class} attains a minimum on the clipped simplex*

$$\Delta_{N,\eta} = \{h \in \mathbb{R}_+^N : h_i \geq \eta, \sum_i h_i = 1\}.$$

If D_s^{class} has Lipschitz gradient on $\Delta_{N,\eta}$, projected gradient descent or cumulative-softmax gradient descent has limit points satisfying the KKT equations (44). If, in addition, the Hessian of D_s^{class} restricted to the tangent space is positive definite at the minimizer, the convergence is locally linear.

Proof. The clipped simplex is compact. Continuity gives existence by Weierstrass. Lipschitz-gradient descent with projection has monotone objective decrease for sufficiently small step size and all accumulation points are first-order stationary. The cumulative-softmax parameterization maps unconstrained logits to the open simplex; after adding a floor $h_i = \eta + (1 - N\eta) \text{softmax}(z)_i$, the same stationarity condition is the pullback of (44). Positive definiteness on the tangent space gives the standard local contraction for gradient descent. $\square$

8 B.8 Worked example: DPM-Solver-v3

For DPM-Solver-v3 in EDM notation, use $\sigma_t = t$, $\alpha_t = 1$, and active EMS coefficients (ℓ_i, s_i, b_i) . The $q=3$ predictor coefficient is (19); the $q=3$ corrector coefficient is (27). The EMS predictor-corrector residual is

$$e_i^{v3} = \gamma_i^P C_{3,i}^{EMS,P} \tilde{g}_i'''(\xi_i) + \gamma_i^C C_{3,i}^{EMS,C} \tilde{g}_i'''(\zeta_i) + R_i^{state}, \quad (51)$$

where

$$C_{3,i}^{EMS,P} = \frac{1}{6} \int_{r_i}^{r_{i+1}} W_i(r; \ell, s, b) (r - r_i)(r - r_{i-1})(r - r_{i-2}) dr, \quad (52)$$

$$C_{3,i}^{EMS,C} = \frac{1}{6} \int_{r_i}^{r_{i+1}} W_i(r; \ell, s, b) (r - r_{i+1})(r - r_i)(r - r_{i-1}) dr. \quad (53)$$

The v3 sequence certificate is

$$D_s^{v3}(\Pi_N) = \sum_i A_i^2 \left| \gamma_i^P C_{3,i}^{EMS,P} + \gamma_i^C C_{3,i}^{EMS,C} \right|^2 V_{3,i}^{EMS}. \quad (54)$$

The implementable grid is

$$\Pi_{N,v3}^* = \arg \min_{0=r_0 < \dots < r_N=1} D_s^{v3}(\Pi_N). \quad (55)$$

The pointwise approximation $m_s(r) \propto d_s(r)^{1/(p+1)}$ is valid only when adjacent gaps satisfy

$$\frac{h_{i-1}}{h_i} = 1 + O(N^{-1}), \quad \frac{h_{i-2}}{h_i} = 1 + O(N^{-1}), \quad (56)$$

and EMS weights vary on an $O(1)$ scale in r . Its leading bias is the first neglected coupling term

$$Bias_{pt}(i) = O \left(A_i^2 V_{3,i} h_i^8 \left[\left| \frac{h_{i-1}}{h_i} - 1 \right| + \left| \frac{h_{i-2}}{h_i} - 1 \right| + \|\partial_r W_i\| h_i \right] \right). \quad (57)$$

At small N , especially $K = 5$, this bias is order-one relative to the leading residual; the sequence optimizer is therefore the certified v3 schedule.

9 B.9 Higher-derivative estimators and pseudocode

For a fine reference grid $\bar{r}_j = j\delta$, $\delta = 1/M$, estimate derivatives of the solver integrand g along a reference trajectory. For $q = 2$, the centered estimator is

$$\hat{g}''(\bar{r}_j) = \frac{g_{j+1} - 2g_j + g_{j-1}}{\delta^2}, \quad Bias = O(\delta^2 \|g^{(4)}\|_\infty). \quad (58)$$

For $q = 3$, use

$$\hat{g}'''(\bar{r}_j) = \frac{g_{j+2} - 2g_{j+1} + 2g_{j-1} - g_{j-2}}{2\delta^3}, \quad Bias = O(\delta^2 \|g^{(5)}\|_\infty). \quad (59)$$

For general q , choose a stencil radius $L \geq \lceil q/2 \rceil$ and coefficients $c_{q,\ell}$ satisfying

$$\sum_{\ell=-L}^L c_{q,\ell} \ell^a = 0 \quad (a < q), \quad \sum_{\ell=-L}^L c_{q,\ell} \ell^q = q!. \quad (60)$$

Then

$$\hat{g}^{(q)}(\bar{r}_j) = \delta^{-q} \sum_{\ell=-L}^L c_{q,\ell} g_{j+\ell}. \quad (61)$$

The variance scales as $\delta^{-2q} \sum_{q,\ell} c_{q,\ell}^2 \text{Var}(g_{j+\ell})$, so increasing M lowers bias but amplifies observational variance.

Listing 1: General- q AB-grid pseudocode.

```
def estimate_gq_squared(net, q, ref_grid, batch, ems=None):
    # Evaluate g or EMS-modulated g_tilde on a fine reference trajectory.
    vals = eval_integrand(net, ref_grid, batch, ems=ems)
    coeffs, radius = finite_difference_stencil(q)
    deriv = apply_stencil(vals, coeffs, radius) / (delta(ref_grid) ** q)
    return mean_square_or_variance(deriv)

def ab_kernel_coef(q, hs, ems_coeffs=None, corrector=False):
    # hs = [h_i, h_{i-1}, ..., h_{i-q+1}]
    # Integrate the Lagrange residual kernel exactly by polynomial algebra
    # or Gaussian quadrature if EMS weights are non-polynomial.
    poly = lagrange_residual_polynomial(q, hs, corrector=corrector)
    if ems_coeffs is None:
        return integrate(poly) / factorial(q)
    W = ems_weight_function(ems_coeffs)
    return integrate(lambda x: W(x) * poly(x)) / factorial(q)

def optimal_abq_grid(table, K, q_map, class_name, max_iter=200):
```

```

# Positive gaps via cumulative softmax.
z = zeros(K)
for it in range(max_iter):
    h = softmax(z)
    r = cumulative([0] + list(h))
    D = 0.0
    for i in range(K):
        q = q_map(i, K)
        hs = history_gaps(h, i, q)
        C = ab_kernel_coef(q, hs, table.ems[i], corrector=table.corrector[i])
        D += table.A[i]**2 * abs(C)**2 * table.Vq[i]
    z = optimizer_step(z, grad(D, z))
return cumulative([0] + list(softmax(z)))

```

One evaluation of D_s^{class} costs $O(Nq_{max})$ when coefficients are pretabulated and $O(Nq_{max}J)$ with J-point quadrature for EMS kernels.

10 B.10 Unified Pareto-domination theorem

Theorem 10.1 (Unified Theorem B: sequence-coupled Pareto domination). *Let $\mathcal{C}$ be a solver class whose leading residual at step i is of the form (33), with order map q_i , terminal amplification A_i , and finite second moments $V_{q_i,i} < \infty$. Assume the grid is monotone, all gaps are positive, EMS weights are bounded and smooth on $[0, 1]$, and the residual kernels $C_{q_i,i}^{class}$ are well-defined. Let*

$$Q_s^{class}(\rho, \Pi_N) = Q_0 + \frac{c_{stat}}{n} \int_0^1 \frac{a(r)}{\rho(r)} dr + c_{disc} D_s^{class}(\Pi_N)$$

with D_s^{class} defined by (41). Define

$$\rho^*(r) = \frac{\sqrt{a(r)}}{\int_0^1 \sqrt{a(z)} dz}, \quad \Pi_{N,s}^* = \arg \min_{0=r_0 < \dots < r_N=1} D_s^{class}(\Pi_N). \quad (62)$$

Then $(\rho^*, \Pi_{N,s}^*)$ minimizes the certified risk Q_s^{class} among all schedules and grids with the same NFE N . Moreover, for every baseline (ρ_b, Π_b) in the same solver class, setting

$$\varepsilon = Q_s^{class}(\rho_b, \Pi_b) \quad (63)$$

and choosing the minimum NFE sequence optimizer satisfying $Q_s^{class} \leq \varepsilon$ yields

$$Q_s^{class}(\rho^*, \Pi^*) \leq Q_s^{class}(\rho_b, \Pi_b), \quad N(\Pi^*) \leq N(\Pi_b). \quad (64)$$

Thus the sequence optimizer weakly Pareto-dominates every non-optimal grid under the certified accuracy-efficiency criterion.

Proof. The statistical term is minimized by Cauchy-Schwarz:

$$\int_0^1 \frac{a(r)}{\rho(r)} dr \geq \left(\int_0^1 \sqrt{a(r)} dr \right)^2,$$

with equality exactly at ρ^* . The discretization term is minimized by definition of $\Pi_{N,s}^*$. Therefore $(\rho^*, \Pi_{N,s}^*)$ minimizes Q_s^{class} at fixed N . For a baseline (ρ_b, Π_b) with N_b , the fixed- N optimum satisfies

$$Q_s^{class}(\rho^*, \Pi_{N_b,s}^*) \leq Q_s^{class}(\rho_b, \Pi_b) = \varepsilon.$$

By definition of the minimum NFE optimizer under the same constraint, $N(\Pi^*) \leq N_b$. The certified risk inequality follows from feasibility. This proves (64). $\square$

Replacement statement for PDF 1 v3. For pointwise-defect solvers, the separable rule $m_s^*(r) \propto d_s(r)^{1/(p+1)}$ is valid. For q-step multistep, predictor-corrector, EMS-modulated, or pseudo-order solvers, replace the scalar pointwise difficulty by the sequence functional $D_s^{class}(\Pi_N)$ in (41), and compute $\Pi_{N,s}^*$ by sequence optimization. The Pareto theorem holds for the corrected sequence optimizer, not for a scalar per-core pointwise $d_s(r)$ rule.

C Third correction: The Finite- N Calibrated Certificate

C.1 Motivation: why the leading-order objective fails at deployment N

PDF 1 v3 proves the leading-order sequence certificate for general multistep, predictor-corrector, EMS-modulated, and pseudo-order solvers. The defect term used there is the Taylor or interpolation asymptotic

$$e_i^{\text{lead}} = C_{s,i}^{q_i}(h_i, h_{i-1}, \dots; l_i, s_i, b_i) g^{(q_i)}(\xi_i), \quad (1)$$

which is valid in the limit $\max_i h_i \rightarrow 0$. At deployment NFE, this leading term may have the correct formal order but the wrong magnitude and ranking. The finite- N correction is to replace the derived asymptotic defect by a measured one-step error surface. The pointwise certificate of Theorem A is already robust for this reason: its $d_s(r)$ is calibrated from measured one-step errors, not analytically derived from a high-order derivative.

The purpose of this appendix is therefore singular: define a calibrated finite- N certificate and prove that Theorem A and Theorem B are its $h \rightarrow 0$ asymptotic shadows.

C.2 Measured local-error surface

Use EDM noise scale $\sigma = t$, half-logSNR $\lambda = -\log \sigma$, and intrinsic coordinate $r \in [0, 1]$ with $r = 0$ at $\sigma_{\max}$ and $r = 1$ at $\sigma_{\min}$. Let

$$\Pi_N = (r_0, \dots, r_N), \quad 0 = r_0 < r_1 < \dots < r_N = 1, \quad h_i = r_{i+1} - r_i. \quad (2)$$

Let $x^{\text{ref}}(r)$ be a high-NFE reference trajectory for the same frozen pretrained backbone, same path family π , and same solver coordinate convention. A q-history solver core s advances one step by

$$x_{i+1}^s = F_s(x^{\text{ref}}(r_i); r_i \rightarrow r_{i+1}, \mathcal{H}_i), \quad (3)$$

where the history $\mathcal{H}_i$ contains the previous reference states and model outputs at

$$r_i, r_{i-1}, \dots, r_{i-q_i+1}. \quad (4)$$

Definition 1 (calibrated one-step error surface). *For solver core s , step order q_i , and gap history $\mathbf{h}_i = (h_i, h_{i-1}, \dots, h_{i-q_i+1})$, define the RMS measured local-error surface*

$$\widehat{e}_s(\sigma_i, \mathbf{h}_i)^2 := \mathbb{E}_{x_0 \sim \mu_0} \|F_s(x^{\text{ref}}(r_i); r_i \rightarrow r_{i+1}, \mathcal{H}_i) - x^{\text{ref}}(r_{i+1})\|_2^2. \quad (5)$$

For $q_i = 1$, the surface reduces to $\widehat{e}_s(\sigma_i, h_i)$. This is the pointwise calibrated difficulty used by Theorem A.

The protocol in (5) is local: start from the reference state, perform exactly one solver step with the specified history, and compare to the reference target. This isolates the finite- N one-step defect from accumulated trajectory drift.

C.3 Calibrated certificate and optimizer

Let $A_i(\Pi_N) \geq 0$ denote terminal amplification from step i to the endpoint. In the calibrated certificate, A_i may be retained as an explicit stability multiplier or absorbed into the measured surface by measuring terminal-propagated one-step errors. The explicit form is

$$D_s^{\text{cal}}(\Pi_N) := \sum_{i=0}^{N-1} A_i(\Pi_N)^2 \widehat{e}_s(\sigma_i, h_i, h_{i-1}, \dots, h_{i-q_i+1})^2. \quad (6)$$

The certified KL/NLL risk becomes

$$Q_s^{\text{cal}}(\rho, \Pi_N) = Q_0 + \frac{c_{\text{stat}}}{n} \int_0^1 \frac{a(r)}{\rho(r)} dr + c_{\text{disc}} D_s^{\text{cal}}(\Pi_N). \quad (7)$$

The calibrated grid is

$$\Pi_{N,s}^{\text{cal},*} \in \arg \min_{0=r_0 < \dots < r_N=1} D_s^{\text{cal}}(\Pi_N). \quad (8)$$

Assumption 1 (regular calibrated surface). *For each admissible history length, $\widehat{e}_s(\sigma, \mathbf{h})$ is continuous on compact subsets of $\sigma \in [\sigma_{\min}, \sigma_{\max}]$ and $\mathbf{h} \in (0, 1]^q$, positive for $h_i > 0$, and strictly increasing in the current gap h_i at fixed $(\sigma, h_{i-1}, \dots)$. It admits a continuous extension to the admissible closed simplex after assigning infinite cost to ill-posed histories with duplicated interpolation nodes.*

Lemma 1 (non-degeneracy of calibrated grids). *Under Assumption 1, the minimizer in (8) exists and has no collapsed adjacent nodes. Moreover, if the leading-order kernel coefficient $C_{s,i}^{q_i}$ has an interior zero for some nonzero gap history, the leading-order objective can admit spurious minimizers that are not minimizers of D_s^{cal} .*

Proof. The admissible ordered simplex is compact after closure. Continuity and the infinite boundary extension give existence and exclude boundary minimizers with duplicated nodes. Strict growth in h_i rules out zero-cost exploitation by shrinking one active step while compensating with an unpenalized history configuration. For the second claim, the leading-order functional contains $|C_{s,i}^{q_i}|^2$. If $C_{s,i}^{q_i} = 0$ at an interior history while $\widehat{e}_s > 0$, the leading-order objective assigns zero asymptotic defect to a step with nonzero measured defect. Hence its minimizer may reduce D_s^{lead} without reducing the true local one-step error. $\square$

C.4 Asymptotic-limit theorem: Theorems A and B as $h \rightarrow 0$ shadows

Let $H_N = \max_i h_i$. For a q-history solver, let the leading kernel coefficient from PDF 1 v3 be

$$C_{s,i}^{q_i} = C_{s,i}^{q_i}(h_i, h_{i-1}, \dots; l_i, s_i, b_i). \quad (9)$$

Let $V_{q_i}(r_i) = \mathbb{E}\|g^{(q_i)}(r_i)\|^2$. If residualization removes the mean defect, replace this by $\text{Var}(g^{(q_i)}(r_i))$.

Assumption 2 (smoothness and nonvanishing leading defect). *The interpolated integrand g has $q_i + 1$ bounded derivatives on each reference step. There are constants $M_{q+1}, \kappa > 0$ such that*

$$\|g^{(q_i+1)}(r)\| \leq M_{q+1}, \quad |C_{s,i}^{q_i}| \sqrt{V_{q_i}(r_i)} \geq \kappa h_i^{q_i+1} \quad (10)$$

for all active steps in the asymptotic regime.

Theorem 1 (asymptotic limit of the calibrated certificate). *Under Assumptions 1 and 10, for multistep, predictor-corrector, EMS-modulated, and pseudo-order solvers,*

$$\boxed{\widehat{e}_s(\sigma_i, h_i, h_{i-1}, \dots)^2 = |C_{s,i}^{q_i}|^2 V_{q_i}(r_i) (1 + O(H_N))}. \quad (11)$$

Consequently,

$$\boxed{D_s^{\text{cal}}(\Pi_N) = D_s^{\text{lead}}(\Pi_N) (1 + O(H_N))}, \quad (12)$$

where

$$D_s^{\text{lead}}(\Pi_N) = \sum_i A_i^2 |C_{s,i}^{q_i}|^2 V_{q_i}(r_i). \quad (13)$$

For a pointwise single-step solver whose effective certificate exponent is p ,

$$\boxed{\widehat{e}_s(\sigma, h)^2 = d_s(\sigma) h^{p+1} (1 + O(h))}, \quad (14)$$

so that

$$\boxed{D_s^{\text{cal}}(m)_o \int_0^1 \frac{d_s(r)}{m(r)^p} dr, \quad m_s^*(r) \propto d_s(r)^{1/(p+1)}}. \quad (15)$$

Thus Theorem A and Theorem B are the two $h \rightarrow 0$ limits of the calibrated certificate (6).

Proof. For a q-history method, the interpolation remainder gives

$$F_s(x^{\text{ref}}(r_i)) - x^{\text{ref}}(r_{i+1}) = C_{s,i}^{q_i} g^{(q_i)}(r_i) + R_{i,q_i+1}, \quad (16)$$

with

$$\|R_{i,q_i+1}\| \leq C_R h_i^{q_i+2} M_{q_i+1}. \quad (17)$$

Squaring and averaging gives

$$\hat{e}_s^2 = |C_{s,i}^{q_i}|^2 V_{q_i}(r_i) + 2\mathbb{E}\langle C_{s,i}^{q_i} g^{(q_i)}, R_{i,q_i+1} \rangle + \mathbb{E}\|R_{i,q_i+1}\|^2. \quad (18)$$

Since $|C_{s,i}^{q_i}| = \Theta(h_i^{q_i+1})$, the cross term is relative order $O(H_N)$ and the last term is relative order $O(H_N^2)$, proving (11). Summing with bounded amplification gives (12). For $q=1$ pointwise methods, the calibrated difficulty is defined by the measured leading RMS-square coefficient d_s , so (14) holds. With local step size $h(r) = 1/m(r)$, the Riemann-sum limit is

$$\sum_i d_s(r_i) h_i^{p+1} o \int_0^1 m(r) d_s(r) m(r)^{-(p+1)} dr = \int_0^1 \frac{d_s(r)}{m(r)^p} dr. \quad (19)$$

Cauchy-Schwarz or the Euler-Lagrange equation gives $m_s^*(r) \propto d_s(r)^{1/(p+1)}$. $\square$

The relative leading-order error is bounded by

$$\left| \frac{D_s^{\text{cal}}}{D_s^{\text{lead}}} - 1 \right| \leq C_s H_N + O(H_N^2), \quad C_s \asymp \frac{M_{q+1}}{\kappa}. \quad (20)$$

Thus a tolerance $\eta \in (0, 1)$ is guaranteed whenever

$$\boxed{H_N \leq h_*(\eta) := \frac{\eta}{2C_s}.} \quad (21)$$

Deployment grids with $H_N > h_*(\eta)$ must use D_s^{cal} , not the leading-order shadow.

C.5 Pareto domination under the calibrated certificate

Theorem 2 (calibrated accuracy-efficiency Pareto domination). *Fix pretrained weights θ , path family π , solver core s , training allocation ρ , and NFE N . Suppose*

$$R_s(\rho, \Pi_N) \leq Q_s^{\text{cal}}(\rho, \Pi_N) \quad (22)$$

with Q_s^{cal} defined by (7). Let $\rho^*(r) \propto \sqrt{a(r)}$ minimize the statistical term and let $\Pi_{N,s}^{\text{cal},*}$ be defined by (8). Then for every baseline grid Π_N^b in the same solver class and same NFE,

$$\boxed{Q_s^{\text{cal}}(\rho^*, \Pi_{N,s}^{\text{cal},*}) \leq Q_s^{\text{cal}}(\rho_b, \Pi_N^b).} \quad (23)$$

Moreover, setting $\varepsilon = Q_s^{\text{cal}}(\rho_b, \Pi_N^b)$, the accuracy-constrained optimizer

$$\Pi_s^*(\varepsilon) \in \arg \min_{\Pi_N} N \quad \text{s.t.} \quad Q_s^{\text{cal}}(\rho^*, \Pi_N) \leq \varepsilon \quad (24)$$

satisfies

$$\boxed{R_s(\rho^*, \Pi_s^*(\varepsilon)) \leq \varepsilon, \quad N(\Pi_s^*(\varepsilon)) \leq N(\Pi_N^b).} \quad (25)$$

Proof. The statistical term is minimized by $\rho^*(r) \propto \sqrt{a(r)}$ by Cauchy-Schwarz. The grid term is minimized by definition of $\Pi_{N,s}^{\text{cal},*}$. Hence (23). If $\varepsilon = Q_s^{\text{cal}}(\rho_b, \Pi_N^b)$, then Π_N^b is feasible for (24). Therefore the minimum-NFE feasible grid has NFE no larger than the baseline. The certified risk bound (22) gives (25). $\square$

C.6 Estimator, sample complexity, and variance advantage

Given a batch $\{x_0^j\}_{j=1}^B$, construct reference trajectories $x_j^{\text{ref}}(r)$ with a high-NFE solver. For each table point $(\sigma, h, h_{-1}, \dots)$, set the reference history, take one step with F_s , and estimate

$$\hat{e}_s^2 = \frac{1}{B} \sum_{j=1}^B \|F_s(x_j^{\text{ref}}(r_i); r_i \rightarrow r_{i+1}, \mathcal{H}_i) - x_j^{\text{ref}}(r_{i+1})\|_2^2. \quad (26)$$

If the reference trajectory has error at most ϵ_{ref} in RMS and the one-step defect is L_F -Lipschitz in the reference state, then

$$\left| \mathbb{E} \widehat{e}_s^2 - \widehat{e}_s^2 \right| \leq C(1 + L_F)^2 \epsilon_{\text{ref}}^2. \quad (27)$$

If the squared one-step defect has variance τ_s^2 , then

$$\text{Var} \left(\widehat{e}_s^2 \right) = \frac{\tau_s^2}{B}. \quad (28)$$

For a table with M_σ noise levels and $M_h^{q_i}$ history-gap configurations, the measurement cost is

$$O(BM_\sigma M_h^{q_i}) \quad (29)$$

one-step solver calls after the reference bank is built.

The leading-order route estimates $g^{(q)}$ by q -th finite differences. If model-evaluation noise has variance σ_g^2 , the q -th finite-difference variance scales as

$$\text{Var}(\widehat{g^{(q)}}) = O(\sigma_g^2 h^{-2q}). \quad (30)$$

The calibrated route estimates a first difference of solver outputs and has variance $O(B^{-1})$ without the factor h^{-2q} . This is the quantitative reason the calibrated pointwise d_s is stable while the derivative-based V_q route becomes unreliable at finite deployment NFE.

C.7 Practical corollary

Build $\widehat{e}_s$ on a grid of $(\sigma, h, h_{-1}, \dots)$ by the one-step reference protocol (5). Interpolate this table during optimization and solve

$$\min_{0=r_0 < \dots < r_N=1} \sum_i A_i^2 \widehat{e}_s(\sigma_i, h_i, h_{i-1}, \dots)^2. \quad (31)$$

A stable parameterization is

$$h_i = \frac{\exp z_i}{\sum_{j=0}^{N-1} \exp z_j}, \quad r_i = \sum_{j < i} h_j. \quad (32)$$

The kernel coefficients $C_{q,i}$, derivative variances V_q , and explicit amplification factors A_i are no longer required as separate analytic constructs if the measured table is built from terminal-propagated local defects. They remain only the $h \rightarrow 0$ explanation of the calibrated certificate.

C.6.1 v4.1 addition: estimator regularity is part of the certificate

Assumption 1 is an assumption on the true calibrated surface

$$\widehat{e}_s(\sigma, h_i, h_{i-1}, \dots, h_{i-q+1}).$$

The implemented optimizer does not minimize this object directly unless the one-step error is evaluated at every candidate grid. It minimizes an estimator or interpolant, denoted

$$\widetilde{e}_s(\sigma, h_i, h_{i-1}, \dots, h_{i-q+1}).$$

Therefore the estimator must preserve the regularity required by Assumption 1. In particular, for each fixed history and noise level it must be nondecreasing in the active gap h_i , and it must not under-report the true surface on the feasible domain used by the optimizer.

Estimator-validity condition. The calibrated-grid optimization in (8) is certified when either

$$(i) \ \widetilde{e}_s = \widehat{e}_s \text{ is evaluated on-the-fly at every candidate grid,}$$

or, for a tabulated estimator, with probability at least $1 - \delta$,

$$\boxed{\widetilde{e}_s(z) \geq \widehat{e}_s(z) \text{ for all feasible } z = (\sigma, h_i, h_{i-1}, \dots), \quad \partial_{h_i} \widetilde{e}_s(z) \geq 0.} \quad (33)$$

A sufficient construction is an upper-envelope interpolant

$$\widetilde{e}_s(z) = I[\widehat{e}_{s,\text{tab}}](z) + c_\delta \text{se}(z) + L_z \text{dist}(z, \mathcal{T}), \quad (34)$$

followed by isotonic projection in h_i , where $\mathcal{T}$ is the measurement table, se is the finite-batch standard error, and L_z is a Lipschitz envelope for the interpolation remainder.

Proposition 3.1 (no blind spots under a valid estimator). *Let $\widetilde{D}_s^{\text{cal}}(\Pi_N)$ be obtained from $D_s^{\text{cal}}(\Pi_N)$ by replacing $\widehat{e}_s$ with an estimator $\widetilde{e}_s$ satisfying (33). Then*

$$D_s^{\text{cal}}(\Pi_N) \leq \widetilde{D}_s^{\text{cal}}(\Pi_N) \quad \text{for every feasible grid } \Pi_N. \quad (35)$$

Hence any risk certificate built from $\widetilde{D}_s^{\text{cal}}$ remains a valid upper certificate for the true calibrated risk. Moreover, if $\widetilde{e}_s$ is strictly increasing in h_i away from the boundary, the non-degeneracy conclusion of Lemma 1 applies to the implemented optimization problem.

Proof. Equation (35) follows termwise from (33):

$$A_i^2 \widehat{e}_s(z_i)^2 \leq A_i^2 \widetilde{e}_s(z_i)^2.$$

Summing over steps gives the claim. Monotonicity in h_i prevents the interpolated surface from creating artificial troughs that the optimizer can exploit; the compactness and boundary-exclusion proof of Lemma 1 then applies verbatim with $\widetilde{e}_s$ in place of $\widehat{e}_s$. $\square$

Failure mode of coarse interpolation. A coarse table can violate (33). If interpolation produces a trough where

$$\widetilde{e}_s(z) < \widehat{e}_s(z),$$

then the optimizer will migrate toward that region because it is minimizing the surrogate surface, not the true calibrated certificate. Such a degenerate grid is not a failure of Theorem 2 or of D_s^{cal} ; it is a failure of the estimator to preserve Assumption 1. Thus the practical rule is: use on-the-fly one-step measurement during grid optimization, or use a table fine enough and conservative enough that its interpolant is an upper, monotone surrogate of the true measured surface.